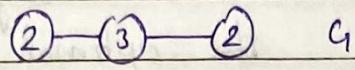


Algorithm Analysis & Design

04/11/2022

Problem set 3

1. (a) Example:

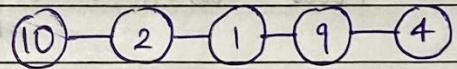


In the above example,

The greedy "heaviest" first algorithm finds
③ to be the max-weight independent
set but ②, ② = 4 is actually the max-
weight independent set.

∴ The heaviest-first greedy approach
does not always give the optimal answer.

- (b) Example:



Our algorithm clubs all odd & even
vertices with each other & chooses the
set with greater weight as MIS.

In the above example, the weight of even
odd sets are

$$\text{even} \rightarrow \{10, 1, 4\} \rightarrow 15$$

$$\text{odd} \rightarrow \{2, 9\} \rightarrow 11$$

so our algorithm says MIS has weight 15.

But we can see that by choosing 10 & 9,
we get MIS as 19.

\therefore This algorithm too, does not yield the optimal solution everytime.

(c)

Let $\text{Opt}(i)$ be the weight of the MIS formed with vertices v_1, v_2, \dots, v_i .

whose weights are w_1, w_2, \dots, w_i respectively.

Now, let S be the set of vertices in MIS of $G_i = \{v_1, v_2, \dots, v_i\}$.

Two cases are possible:

Case 1: $v_i \in S$

If $v_i \in S$, then $v_{i-1} \notin S$ since two adjacent vertices are in S .

$$\therefore \text{Opt}(i) = \text{Opt}(i-2) + w_i$$

Case 2: $v_i \notin S$

If $v_i \notin S$, then $\text{Opt}(i)$ will simply be equal to

$$\text{Opt}(i) = \text{Opt}(i-1)$$

Since we do not know whether v_i is present or not, we consider both of the cases and take the max of the two.

Doing so, we are considering all possible independent sets and picking the one with max value.

$$\therefore \text{Opt}(i) = \max \left\{ \begin{array}{l} \text{opt}(i-2) + w_i \\ \text{opt}(i-1) \end{array} \right\}$$

with $\text{opt}(0) = 0$ & $\text{opt}(1) = w_1$, as the base condition.

Algorithm:

We'll walk through the path from left to right, using the memoized minimum weight so far to decide which which property must apply (either the last vertex is a part of the set or not)

Initialize Array A of size $n+1$

$$A[0] = 0, A[1] = w_1$$

for i in $n-1$

$$A[i+1] = \max(A[i], w_i + A[i-1])$$

Set $S \leftarrow \emptyset$.

$i = n$

While $i \geq 1$:

if $A[i] = A[i-1]$ # last vertex not in set

$\not\in i--$

else # last vertex in set

$S.insert(i)$

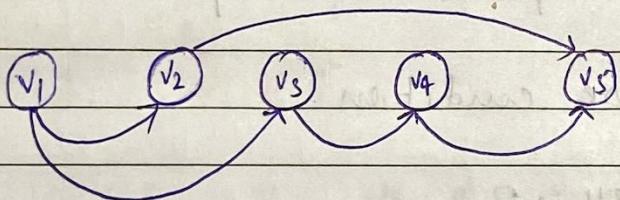
$i -= 2$

end

end

return S ;

2. (a)



The algorithm will give the length of the longest path as 2, as it will choose the path $v_1 \rightarrow v_2 \rightarrow v_5$, but there is a longer path, $v_1 \rightarrow v_3 \rightarrow v_4 \rightarrow v_5$ of length 3 which is the longest.

∴ The specified algorithm does not correctly solve the problem.

(b)

(b) let $\text{opt}(i)$ be the length of the longest path from v_i to v_j .

and let $\text{opt}(1) = 0$.

If there is no path from v_i to v_j ,

set $\text{opt}(i) = -\infty$

Now, $\text{opt}(i)$ can be written as follows:

$$\text{opt}(i) = 1 + \max_{1 \leq j < i} (\text{opt}(j)) \text{ if } \exists \text{ path from } v_j \text{ to } v_i.$$

i.e. we choose the vertex which has

the longest path from v_j just preceding i which also has path from v_j to v_i .

∴ we now need to find $\text{opt}(n)$.

We can do this iteratively by creating an array & building the solution from 1 to n

Algorithm / Pseudo code:

Longest-path(n) :

$A[1, \dots, n] \leftarrow -\infty$

$A[0] = 0$

for $i = m \dots n$

if $A[i] \neq -\infty \text{ & } S[i] \leq A[i] + 1$

for $i \in m \dots n$

for all edges (j, i) :

if $M[j] \neq -\infty \text{ & } M[i] < M[j] + 1$:

$$M[i] = M[j] + 1$$

Return $M[n]$.

The run-time is $O(n^2)$ as the

max. number of edges can be $1 + 2 + \dots + n-1$

$$= \frac{n(n-1)}{2} \approx O(n^2)$$

and assuming all of them can/must be traversed

3. Let $\text{opt}(i)$ denote the min. cost
of a solution till week i .

For week i either company A or
B will be picked.

Case 1: Company A is picked.

$$\text{then } \text{opt}(i) = r_C + \text{opt}(i-1)$$

and Case 2: Company B is picked

$$\text{then } \text{opt}(i) = r_B + \text{opt}(i-4)$$

we do not know which of the two is minimum, therefore we choose the minimum of the two.

$$\therefore \text{opt}(i) = \min(r_{S_i^*} + \text{opt}(i-1), t_c + \text{opt}(i-1)).$$

for $i \leq 0$, let $\text{opt}(i) = 0$.

$\therefore \text{opt}(n)$ can be found by building up the opt values in increasing order of i assuming each iteration takes constant time.

We can obtain the optimal schedule by tracing back through the arrays of opt values.

4. Let $\text{opt}(i)$ be the total cost of the optimal configuration from server 1 to i , given there is a copy of the file at server j .

Let j^* be the place where the copy of the file is located, which is an immediate predecessor of server j .

The total cost of the optimal configuration, $\text{opt}(i)$, can be written as follows:

$$\text{opt}(i) = \text{placement cost at } j^* + \text{access time cost from } j^* \text{ to } i.$$

placement cost at $i - c_i$

and access time cost from $j+1$ to i .

$$= 0 + 1 + 2 + \dots + i-j = \frac{(i-j)(i-j-1)}{2}$$

$\downarrow \quad \downarrow \quad \downarrow \quad \downarrow$
 $\text{for } i \quad \text{for } i-1 \quad \text{for } i-2 \quad \text{for } j+1$

We do not know where j might be,

$$\therefore \text{opt}(i) = c_g + \min_{0 \leq j < i} (\text{opt}(j) + \frac{(i-j)(i-j-1)}{2})$$

with $\text{opt}(0) = 0$.

The values of opt can be built in a bottom up manner in $\text{O}(i)$ for i^{th} iteration

leading to a total run time of $\text{O}(n^2)$.

We can also find the optimal configuration by tracing through the values of opt in the array that stores it.

5.

The following algorithm checks the validity of the given $d(v)$'s in $O(m)$ time:

If for any edge $e = (v, w)$, we have

$$d(v) > d(w) + c_e$$

we can immediately reject.

This follows since if $d(w)$ is correct, then there is a path from v to t via w of cost $d(w) + c_e$. The minimum distance from v to t is at most this value, so $d(v)$ must be incorrect.

Now, consider the graph G' formed by taking G and removing all edges except $e = (v, w)$ for which $d(v) = d(w) + c_e$.

We will now check that every node has a path to t in this new graph.

This can be easily checked in $O(m)$ by reversing all edges and doing a DFS or BFS.

If any nodes fails to have such a path, reject it and accept otherwise.

observe that if $d(v)$ were correct for all nodes v , then if we consider those edges on the shortest path from any node v to t , these edges will all be in G' . Therefore we can safely reject if any node can not reach t in G' .

So far, we have argued that when our algorithm rejects a set of distances, those distances must be incorrect.

We now need to prove that if our algorithm accepts, then the distances are correct.

Let $d'(v)$ be the actual distance in G from v to t . We need to show $d'(v) = d(v) + v$.

Consider the path found by our algorithm in G' from v to t .

The real cost of this path is exactly $d(v)$. There may be shorter path, but we know that $d'(v) \leq d(v)$, and this holds $\forall v$.

Now suppose that there is some v for which $d'(v) < d(v)$.

Consider the actual shortest path P from v to t . Let n be the last on P for which $d'(n) < d(n)$. Let y be the node after n on P . By our choice of n ,

$$d'(y) = d(y)$$

Since P is the real shortest path to t from v , it is also the real shortest path from all nodes on P .

$$\text{So } d'(n) = d'(y) + c_e \text{ where } e = (n, y)$$

This implies that $d(n) > d'(n) + c_e$. \dagger

But then we have a contradiction, since our algorithm would have rejected upon reject inspection of e .

$$\therefore d'(v) = d(v) \neq v.$$

The same solution can be phrased in terms of the modified cost function that will be given in part (b) of the solution.

There are two common mistakes:

- 1) To simply use the algorithm that just checks for $d(v) > d(w) + c_e$.

This can be broken if the graph has a cycle of cost 0, as the nodes on such a cycle may be self-consistent, but may have a much larger distance to the root than reported.

Consider two nodes n & y connected to each other with 0 cost edges, and connected to t by an edge of cost 5.

If we report that both are at distance 3 from the root t , the algorithm will ~~fail~~ wrongly accept the distances.

2) To fail to prove the algorithm output is correct on both yes & no instances of the problem.

(b) Given distances to a sink t , we can efficiently calculate distance to another sink t' in $O(m\log n)$ time.

To do so, note that if only edge costs were non-negative, then we could just run Dijkstra's algorithm.

We will modify costs to ensure this, without changing the min cost paths.

Consider modifying the cost of each edge $e = (v, w)$ as follows:

$$c'_e = c_e - d(v) + d(w)$$

First observe that the modified costs are non-negative, since if $c'_e < 0$ then $d(v) < d(w) + c_e$ which is not possible if d reflect true distances to t .

Now, consider any path P from some node v to t' . The real cost of P is

$$c(P) = \sum_{e \in P} c_e$$

The modified cost of P is $c'(P)$

$$c'(P) = \sum_{e \in P} c'_e = \sum_{e \in P} c_e - d(v) + d(w)$$

Note that all but the first and last node in P occur once positively & once negatively in this sum, so we get that $c'(P) = c(P) - d(v) + d(w)$.

Hence the modified cost of any path from v to w differs from the real cost of that same path an additive

$d(t') - d(v)$, a constant. So the set of minimum cost paths from v to w under c' is the same as the set under c .

Furthermore, given the min. distance from v to t' under c , we can calculate the minimum distance under c' by simply adding $d(v) - d(t')$.

Our algorithm is exactly that:

- 1) Calculate the modified costs
- 2) Run Dijkstra from t'
- 3) Adjust the distances as described.

The edges need to be reversed for standard implementation of Dijkstra's

The time complexity of the algorithm is

$$\Theta(m + m \log n + n) = \Theta(m \log n) \text{ time.}$$