Today: Hubs, Switches, and Routers, Oh My!

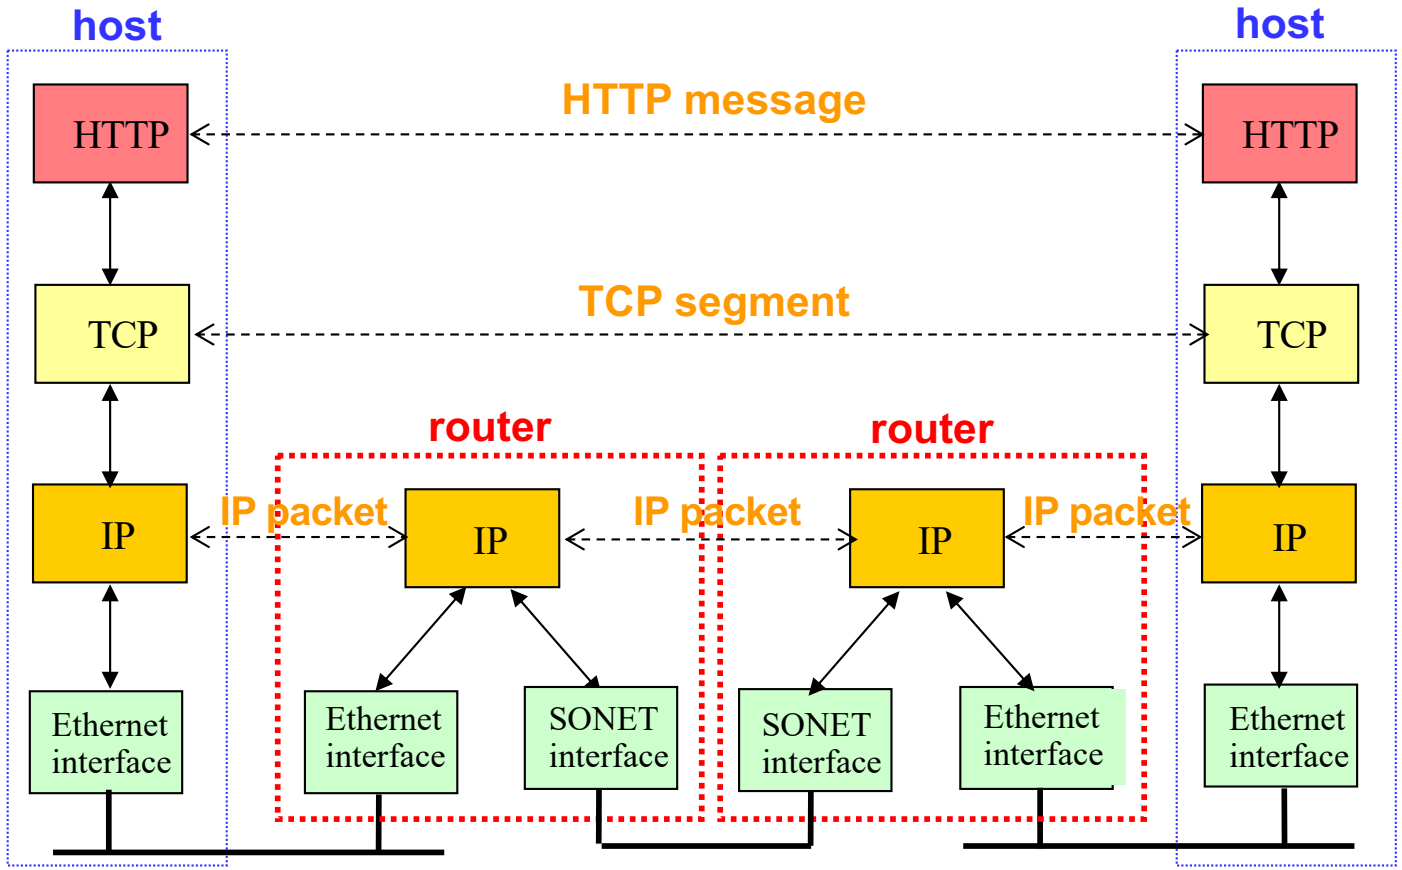

Terminology

- Hubs and Repeaters
  - Connect machines on same “layer 2” LAN
  - Broadcast: All frames are sent out all physical ports
- Switches and Bridges
  - Connect machines on same “layer 2” LAN
  - Only send frames to selected physical port based on destination MAC address
- Routers
  - Connect between LANs at “layer 3”, e.g., wide area
  - Only send packet to selected physical port based on destination IP address

“Layer 2”
Hubs and Switches

Physical Layer: Repeaters

- Distance limitation in local-area networks
  - Electrical signal becomes weaker as it travels
  - Imposes a limit on the length of a LAN
- Repeaters join LANs together
  - Analog electronic device
  - Continuously monitors electrical signals
  - Transmits an amplified copy

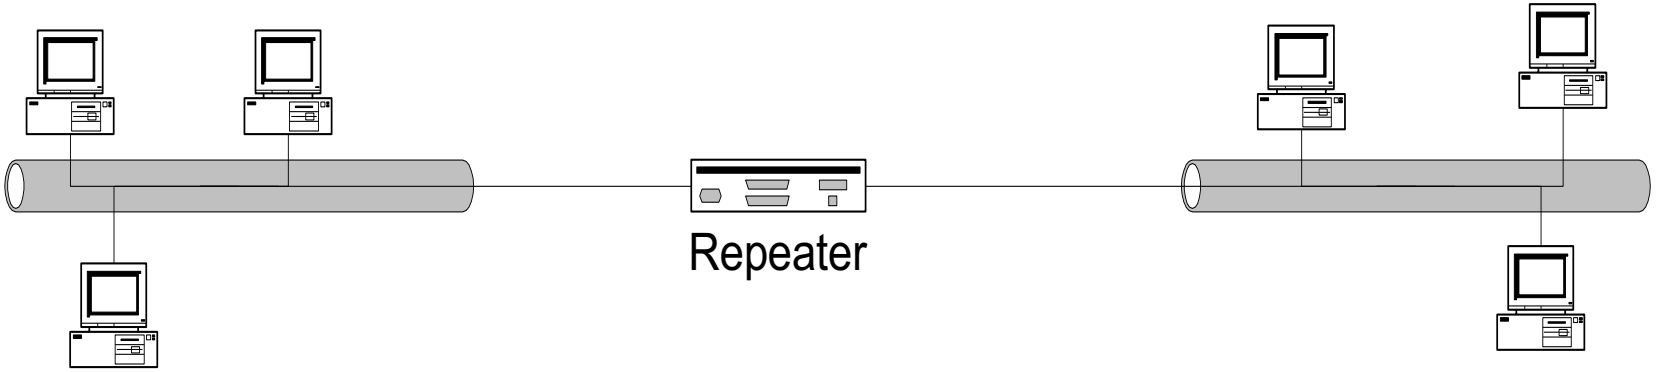

Physical Layer: Hubs

- Joins multiple input lines electrically
  - Designed to hold multiple line cards
  - Do not necessarily amplify the signal
- Very similar to repeaters
  - Also operates at the physical layer

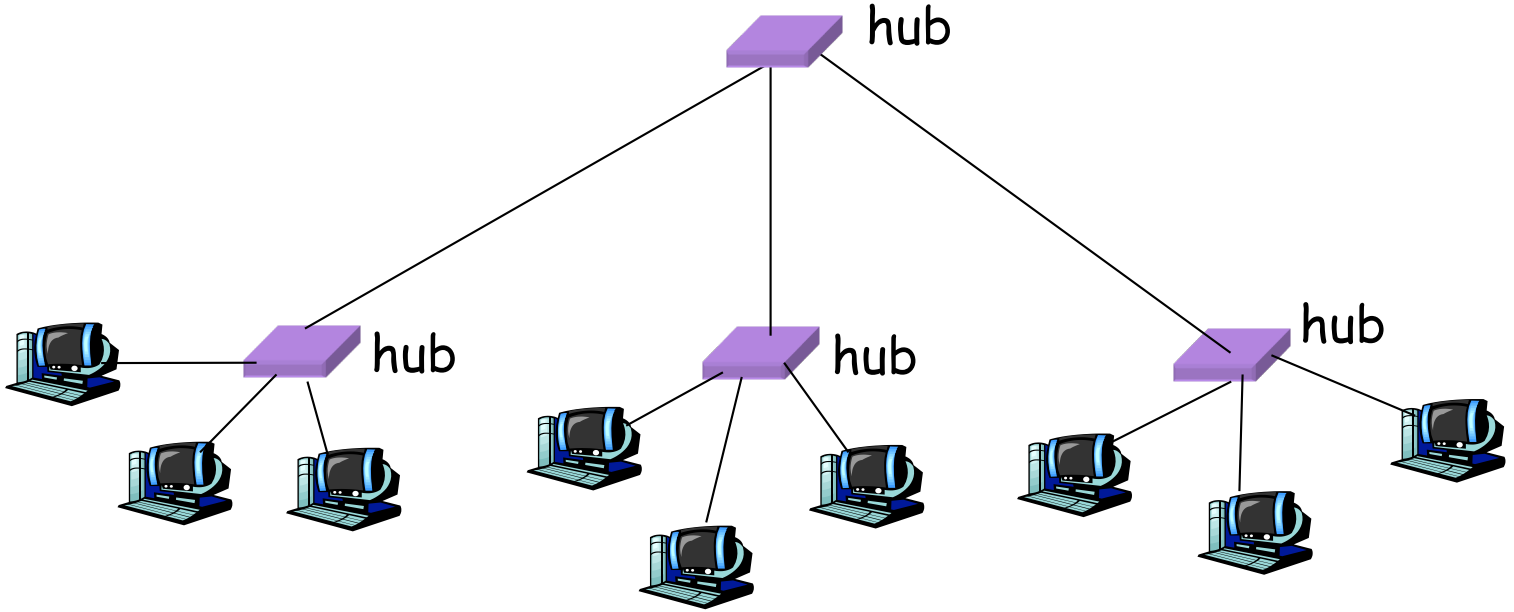

Hub: Overview

- **Layer 1** Device
- **1** Collision Domain
- **Half-**Duplex
- **Wasted** Bandwidth
- **Security** Risks
- **Replaced** by Switches

Limitations of Repeaters and Hubs

- One large shared link
  - Each bit is sent everywhere
  - So, aggregate throughput is limited
- Cannot support multiple LAN technologies
  - Does not buffer or interpret frames
  - Can't interconnect between different rates/formats
- Limitations on maximum nodes and distances
  - Shared medium imposes length limits
  - E.g., cannot go beyond 2500 meters on Ethernet

Link Layer: Bridges

- Connects two or more LANs at the link layer
  - Extracts destination address from the frame
  - Looks up the destination in a table
  - Forwards the frame to the appropriate segment
- Each segment can carry its own traffic

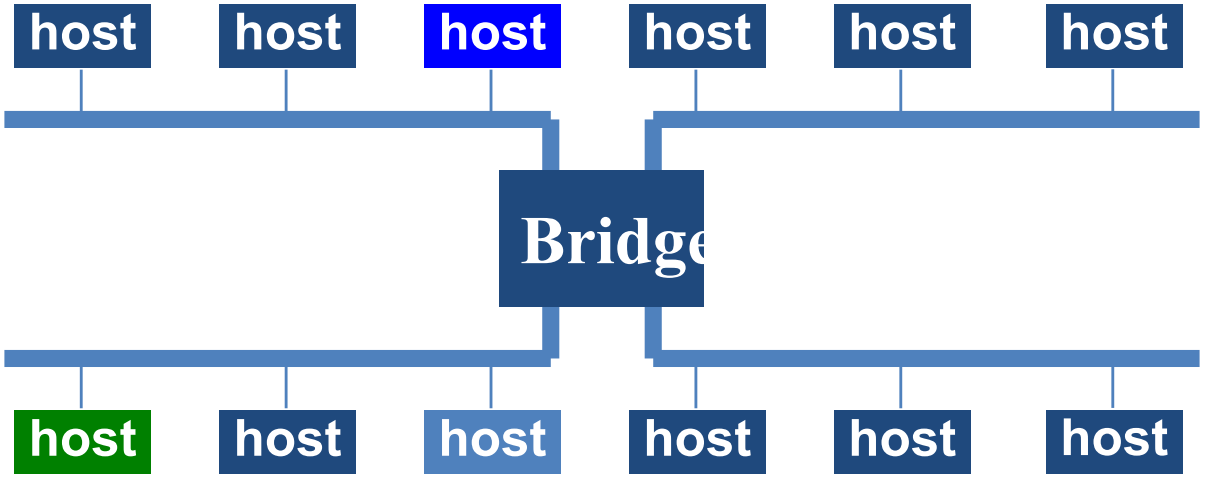

Bridge: Overview

- **Layer 2** Device
- **Segments** Lans
- **2** Collision Domains
- **Fewer** Ports
- **Replaced** by Switches

Link Layer: Switches

- Typically connects individual computers
  - A switch is essentially the same as a bridge
  - … though typically used to connect hosts
- Supports concurrent communication
  - Host A can talk to C, while B talks to D

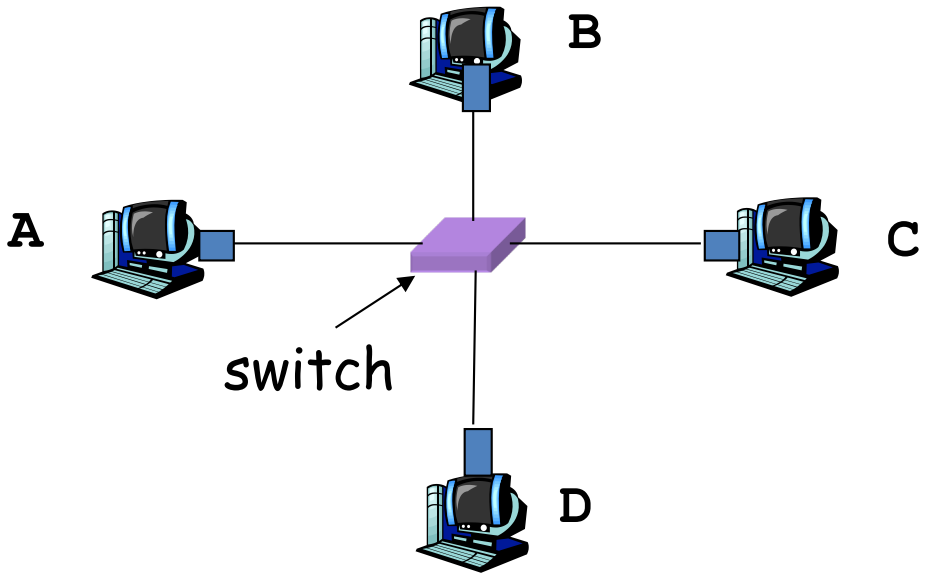

Switch: Overview

- **Layer 2** Device
- **Full-**Duplex
- **Multiple** Collision Domains
- **Saves Bandwidth**
- **Increased Security**

Bridges/Switches: Traffic Isolation

- Switch filters packets
  - Frame only forwarded to the necessary segments
  - Segments can support separate transmissions

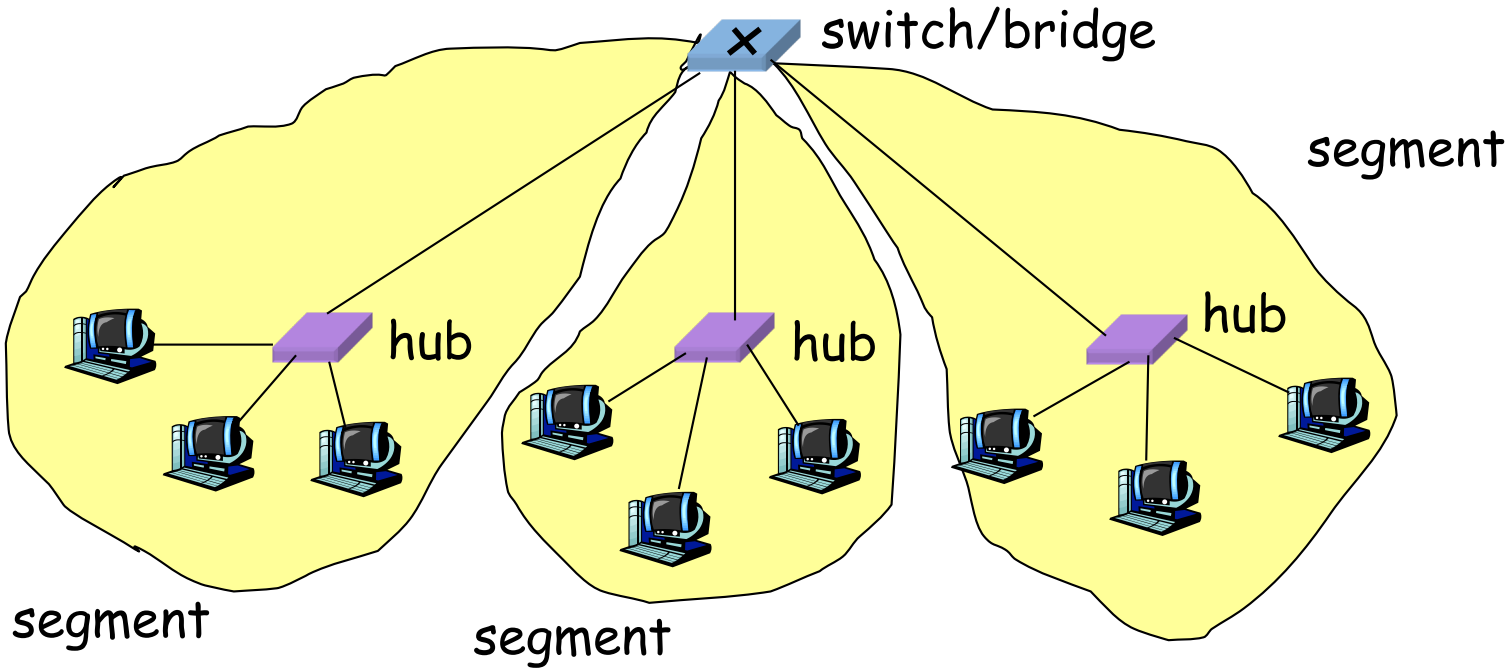

Self Learning: Building the Table

- When a frame arrives
  - Inspect the *source* MAC address
  - Associate the address with the *incoming* interface
  - Store the mapping in the switch table
  - Use a timer to eventually forget the mapping

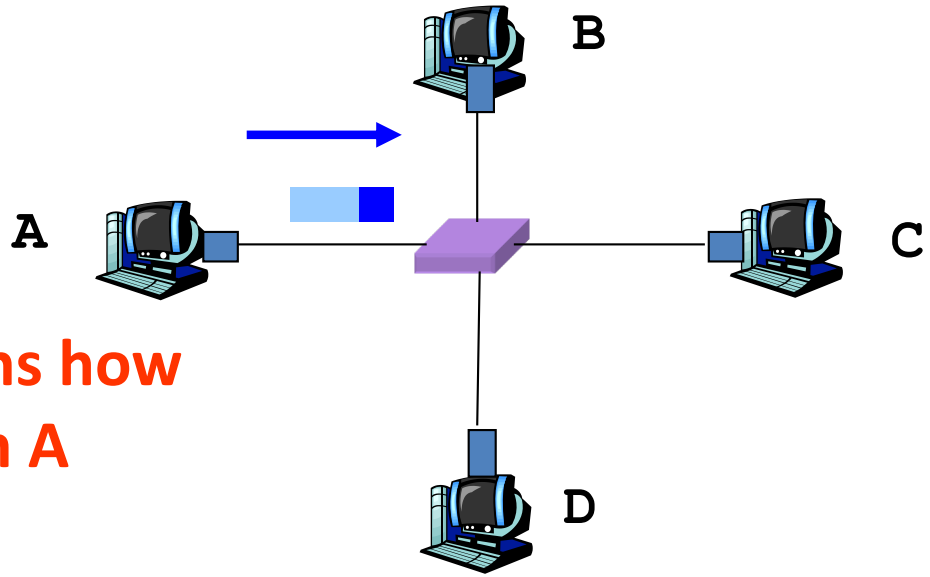

Switch learns how to reach A

Self Learning: Handling Misses

- When frame arrives with unfamiliar destination
  - Forward the frame out all of the interfaces
  - … except for the one where the frame arrived
  - Hopefully, this case won't happen very often!

When in doubt, shout!

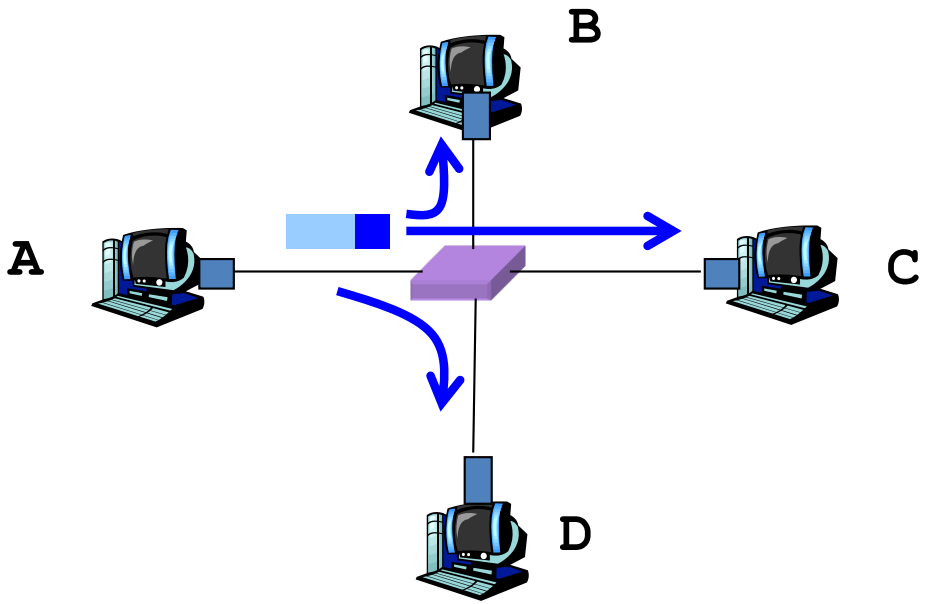

Switches vs. Hubs

- Compared to hubs, Ethernet switches support
  - (Y) Larger geographic span
  - (M) Similar span
  - (C) Smaller span
- Compared to hubs, switches provide
  - (Y) Higher load on links
  - (M) Less privacy
  - (C) Traffic isolation

Routers: Looking closer…

Router: Overview

- **Routes Traffic** Between Networks
- **Layer 3** Device
- **Fewer Ports**

Basic Router Architecture

- Each switch/router has a forwarding table
  - Maps destination address to outgoing interface

- Basic operation
  1. Receive packet
  2. Look at header to determine destination address
  3. Look in forwarding table to determine output interface
  4. Modify packet header (e.g., decr TTL, update chksum)
  5. Send packet to output interface

Basic Router Architecture

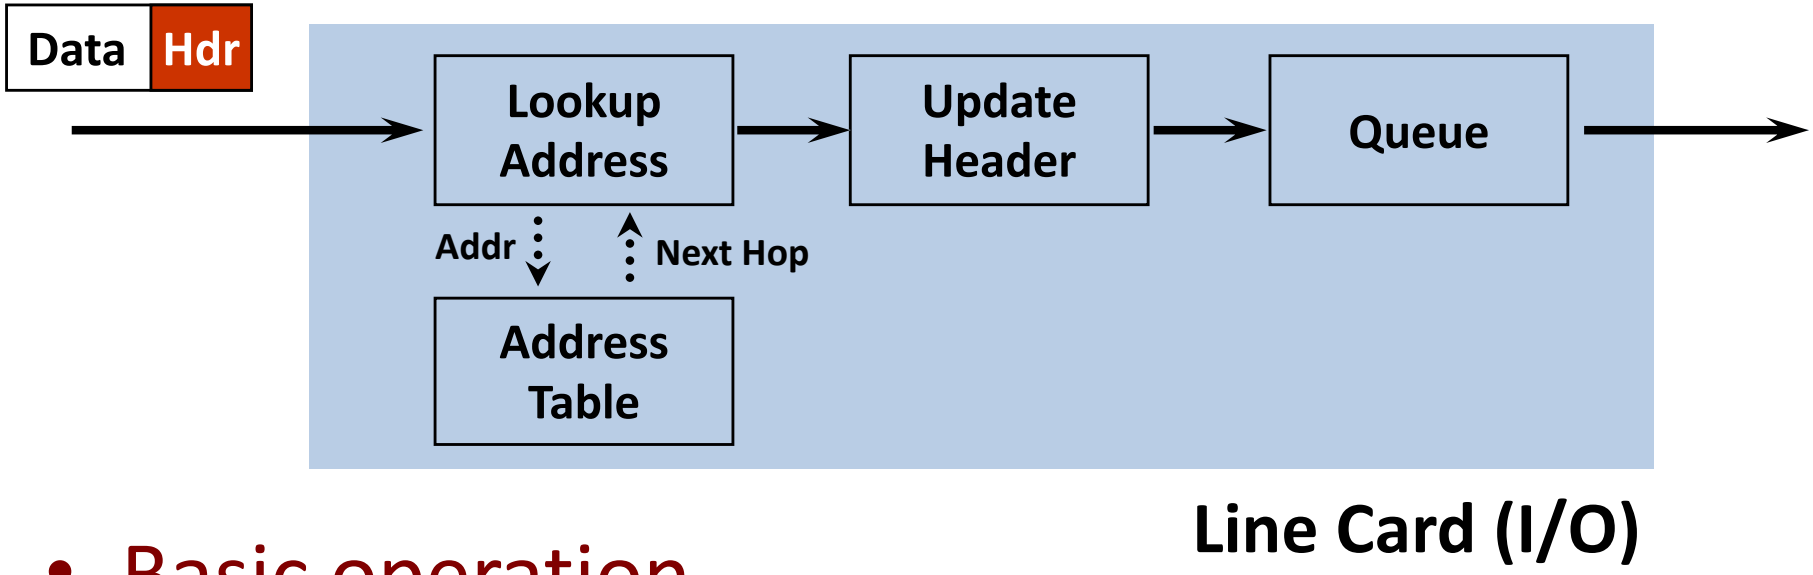

- Basic operation
  1. Receive packet
  2. Look at header to determine destination address
  3. Look in forwarding table to determine output interface
  4. Modify packet header (e.g., decr TTL, update chksum)
  5. Send packet to output interface

Router

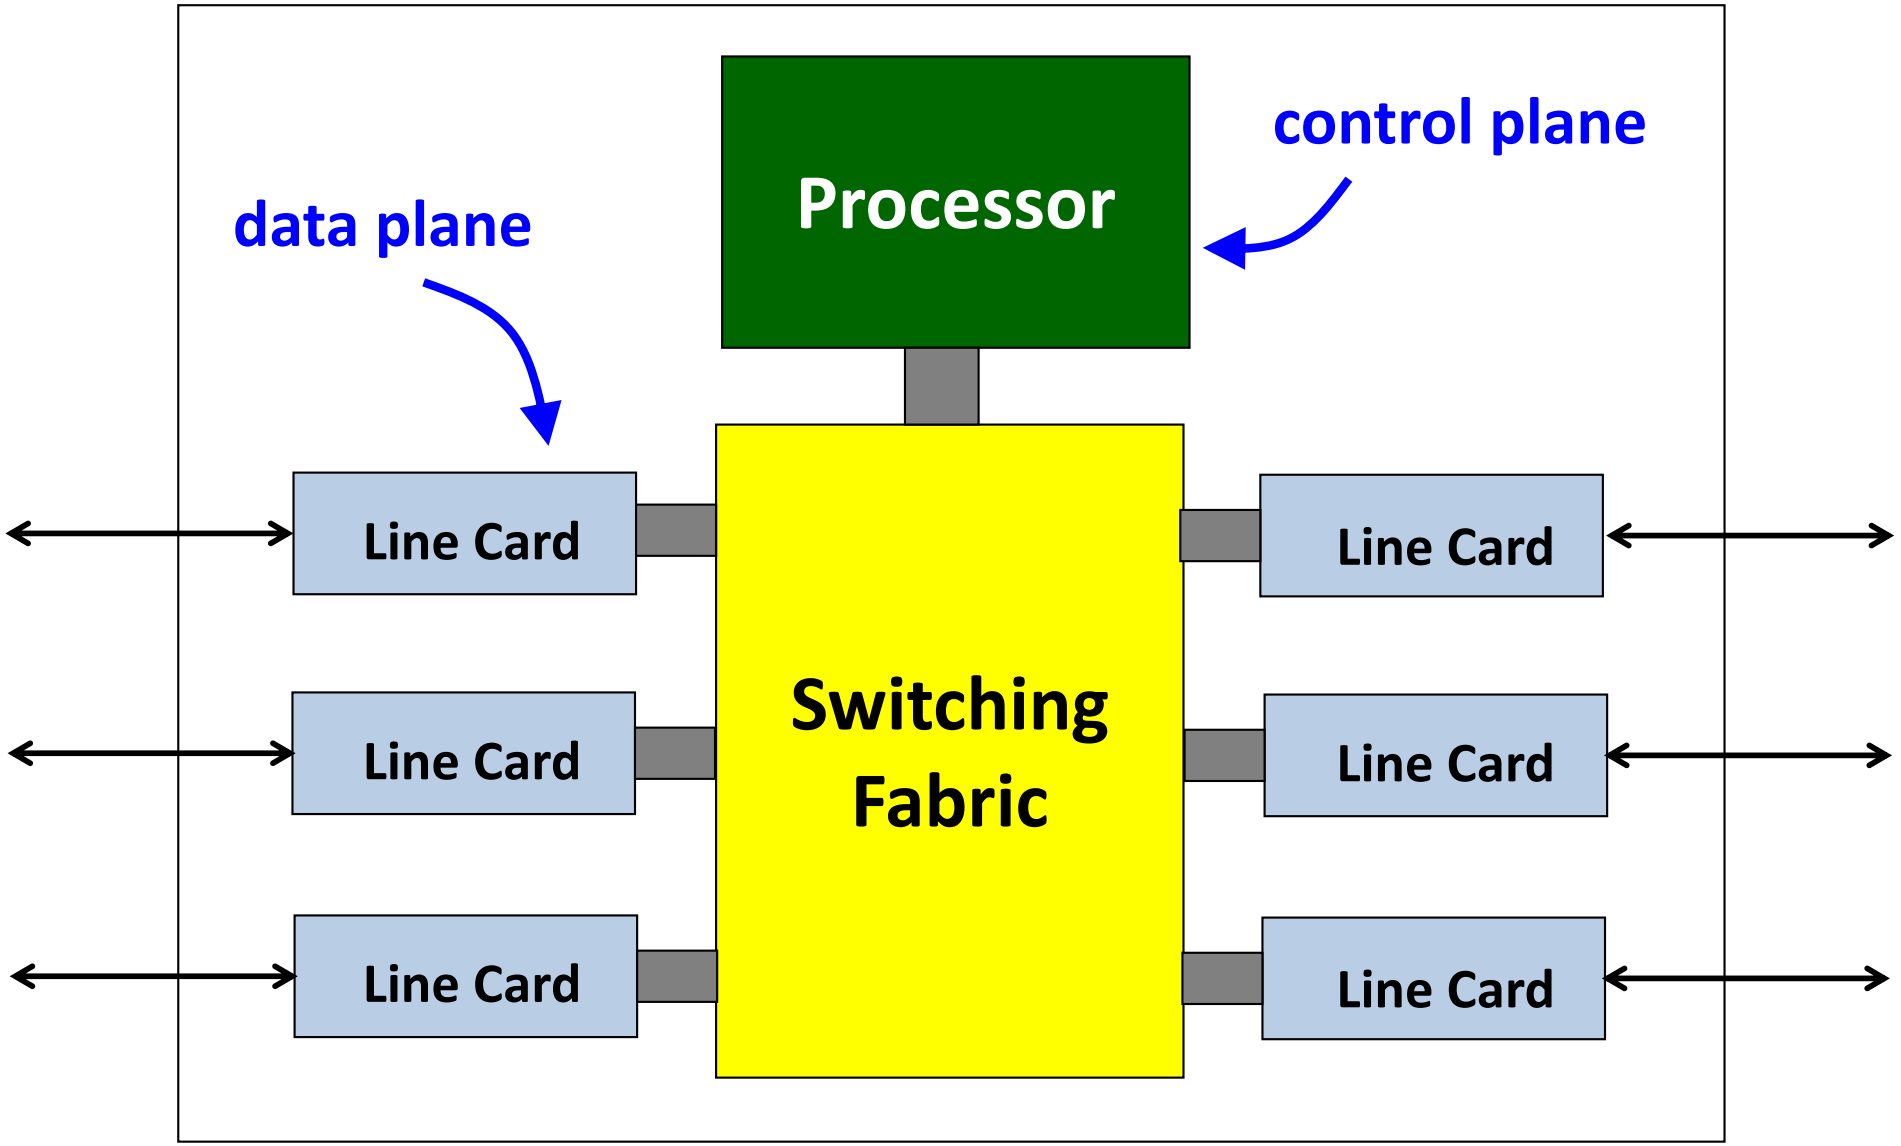
Processor
control plane
data plane
Line Card
Line Card
Line Card
Switching Fabric
Line Card
Line Card
Line Card

Router

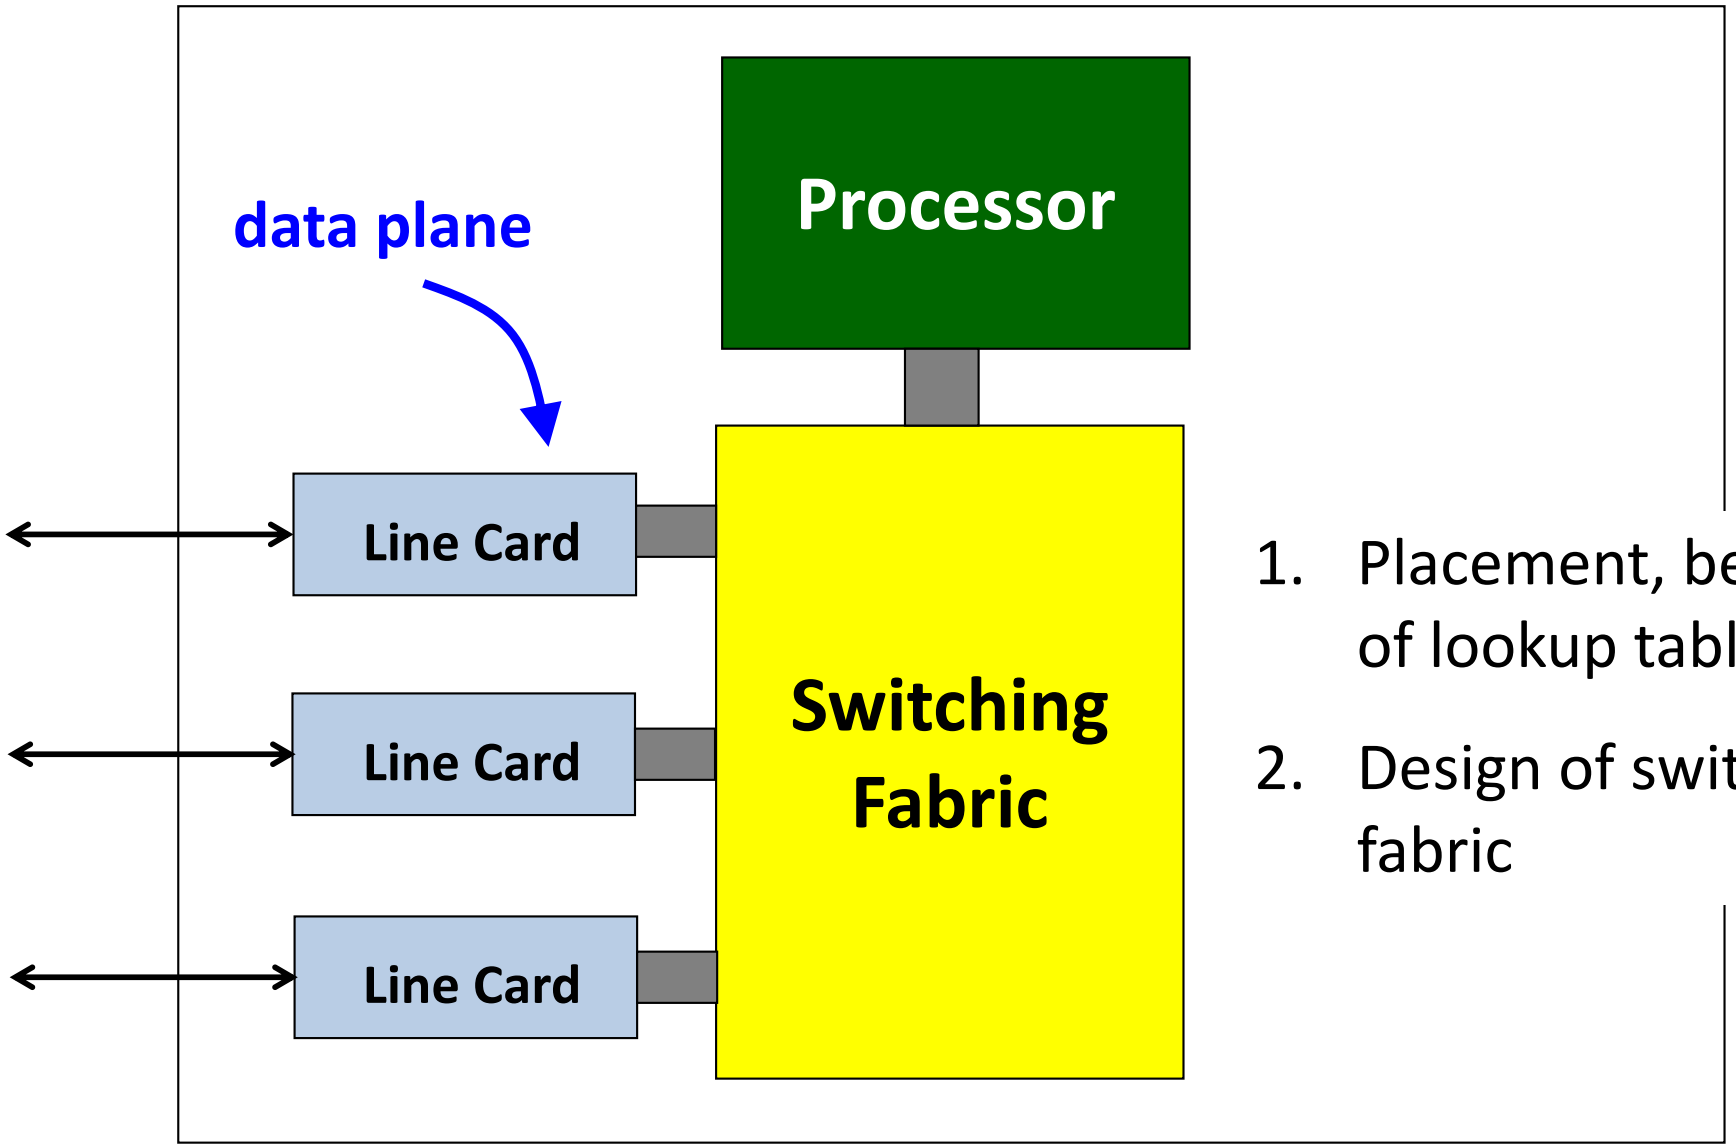

1. Placement, behavior of lookup tables
2. Design of switching fabric

Lookup algorithm depends on protocol

Protocol	Mechanism	Techniques
Ethernet (48 bits) MPLS ATM	Exact Match	• Direct lookup • Associative lookup • Hashing • Binary tree
IPv4 (32 bits) IPv6 (128 bits)	Longest-Prefix Match	• Radix trie • Compressed trie • TCAM

Longest Prefix Match (LPM)

- Each packet has destination IP address
- Router looks up table entry that matches address

68.211.6.120

Prefix	Output
68.208.0.0/12	1
68.211.0.0/17	1
68.211.128.0/19	2
68.211.160.0/19	2
68.211.192.0/18	1

LPM: Motivation

- Each packet has destination IP address
- Router looks up table entry that matches address

- Benefits of CIDR allocation and LPM
  - **Efficiency:** Prefixes can be allocated at much finer granularity
  - **Hierarchical aggregation:** Upstream ISP can aggregate 2 contiguous prefixes from downstream ISPs to shorter prefix

Software LPM lookup using trie

- Prefixes “spelled out” by following path from root
- To find the best prefix spell out address in trie

	Prefixes
A	111*
B	10*
C	1010*
D	10101

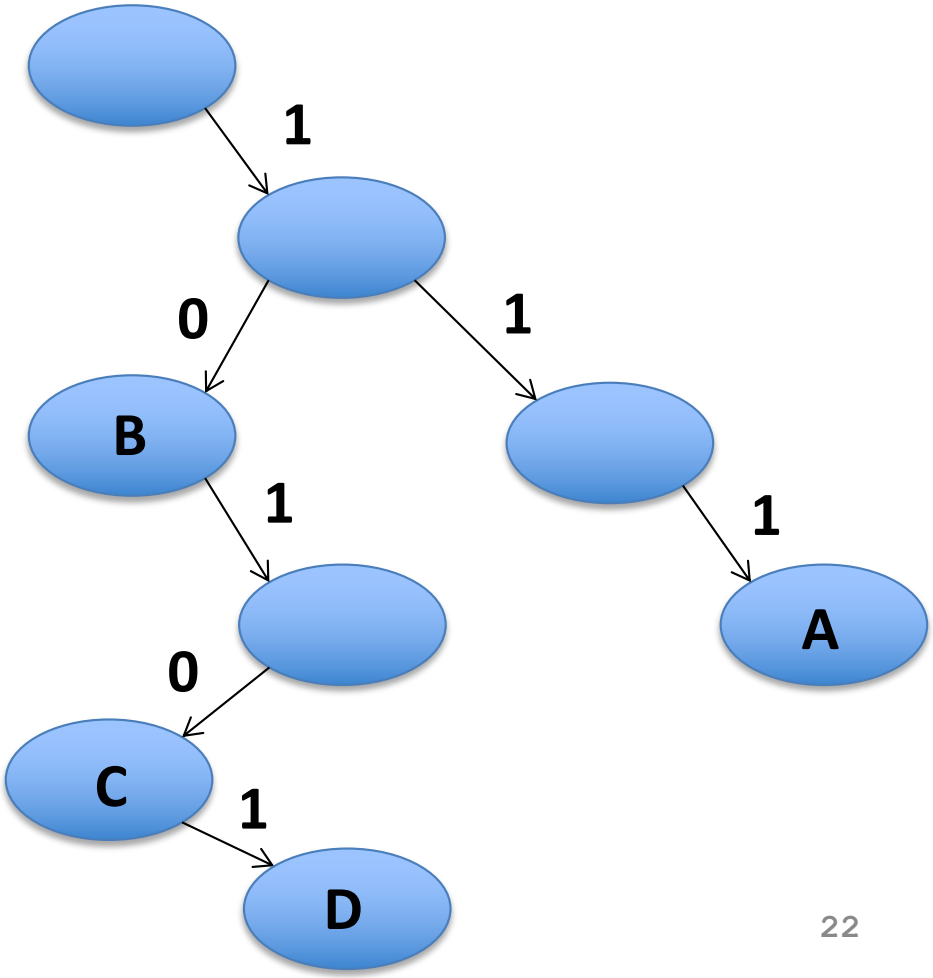

Software LPM lookup using trie

- Prefixes "spelled out" by following path from root
- To find the best prefix spell out address in trie

- 1 lookup per level → max 32 lookups/address!
- Too slow:
  - E.g., "Optical Carrier 48" line (2.5 Gbps) requires 160ns lookup … or 4 memory accesses

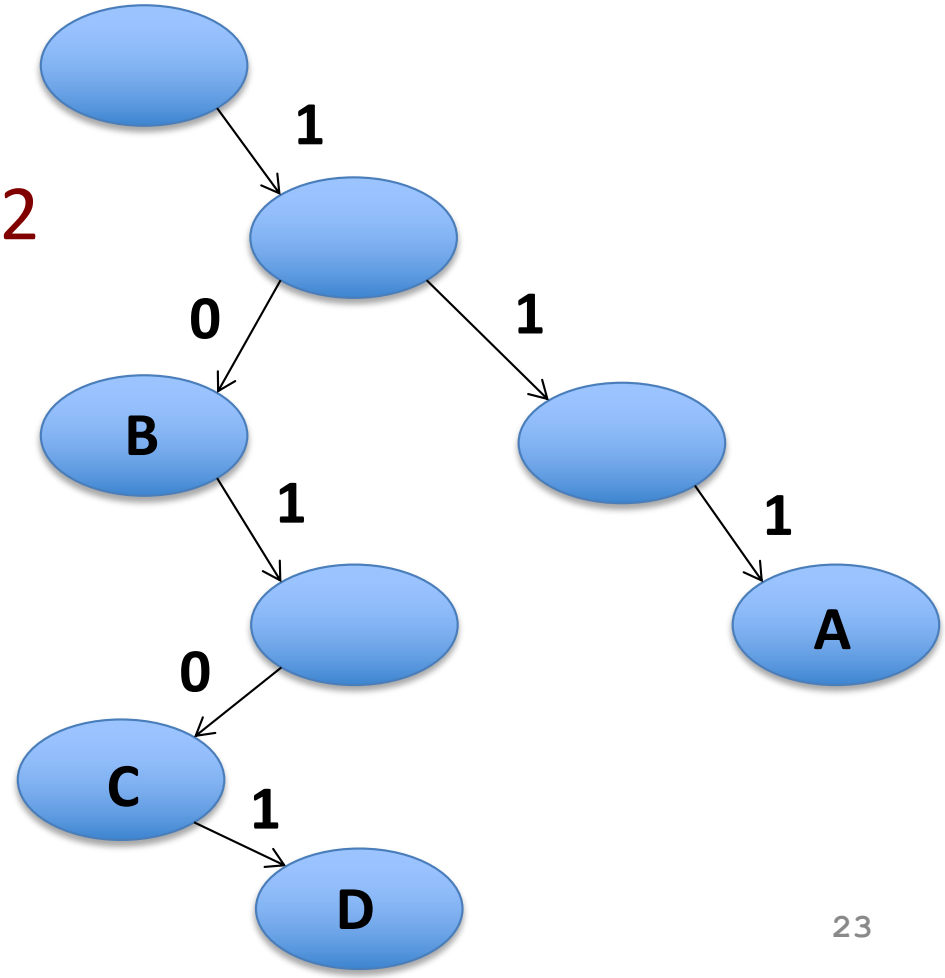

Software LPM lookup: k-ary trie (k=2)

	Prefixes
A	111*
B	10*
C	1010*
D	10101

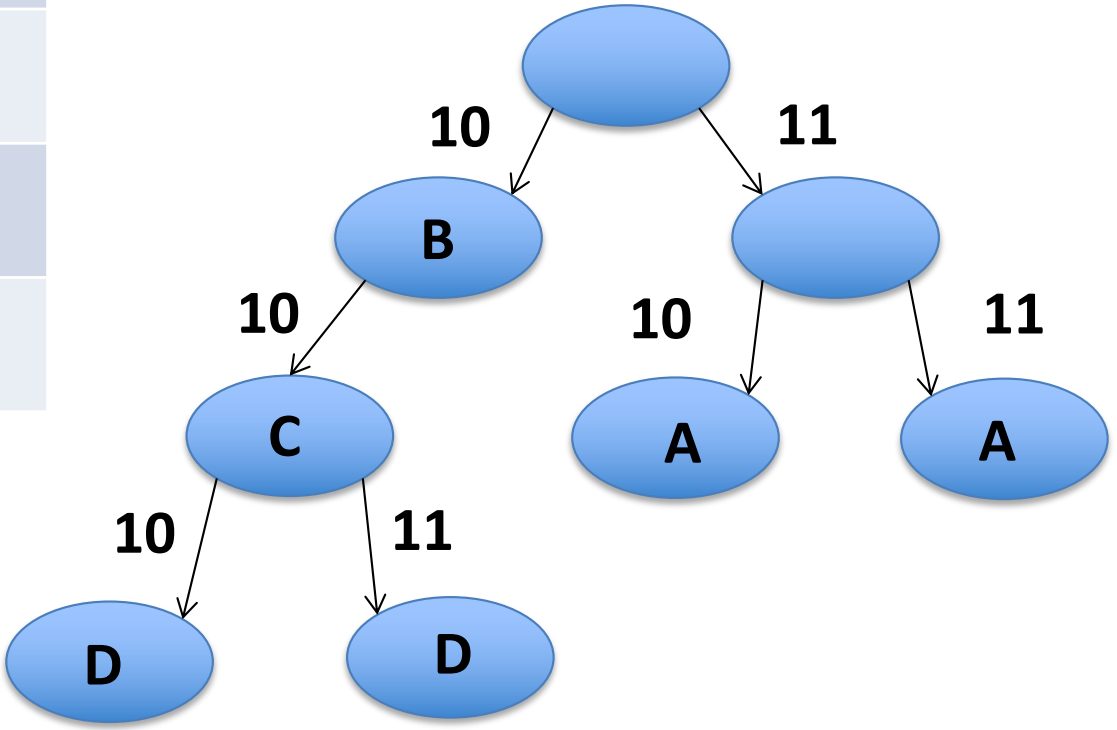

IP Address Lookup: Binary Tries

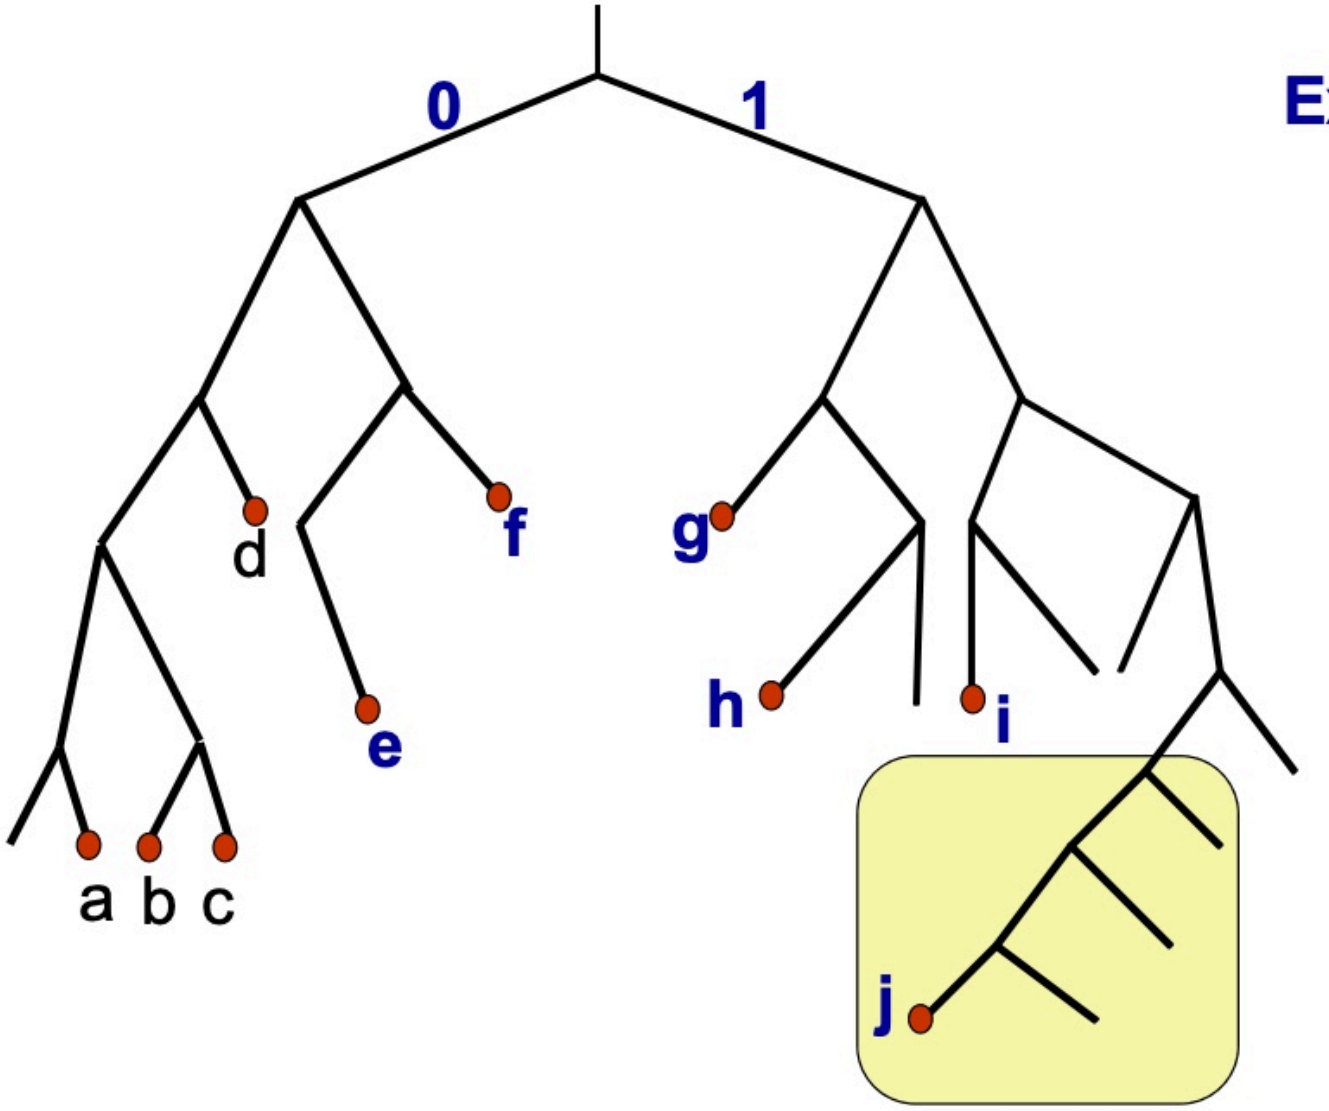

Example Prefixes:

a) 00001
b) 00010
c) 00011
d) 001
e) 0101
f) 011
g) 100
h) 1010
i) 1100
j) 11110000

IP Address Lookup: Patricia Trie

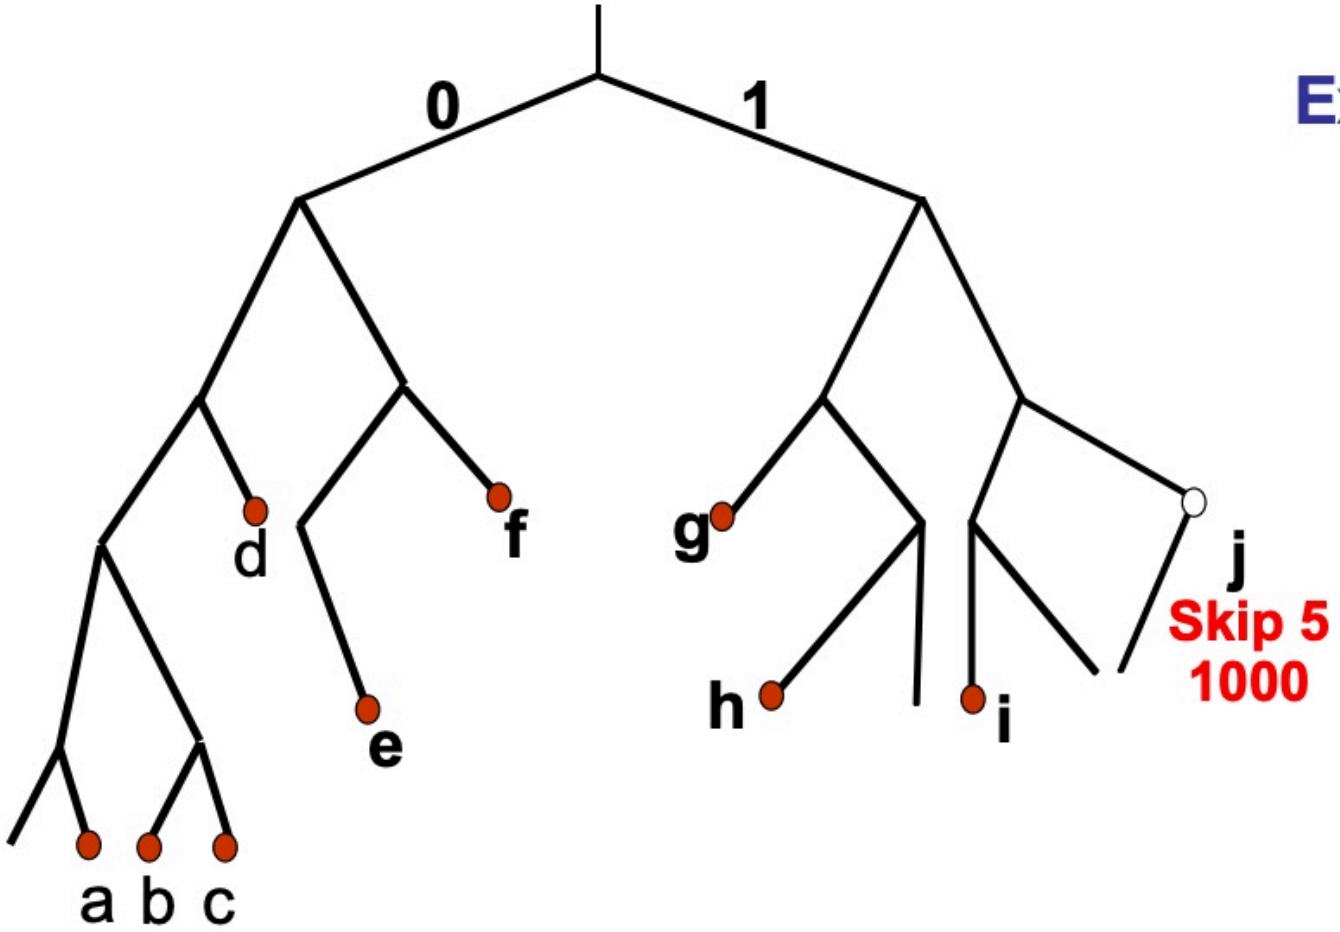

Example Prefixes

a) 00001
b) 00010
c) 00011
d) 001
e) 0101
f) 011
g) 100
h) 1010
i) 1100
j) 11110000

Problem: Lots of (slow) memory lookups

LPM with PATRICIA Tries

- Traditional method – Patricia Tree
  - Arrange route entries into a series of bit tests
- Worst case = 32 bit tests
  - Problem: memory speed, even w/SRAM!

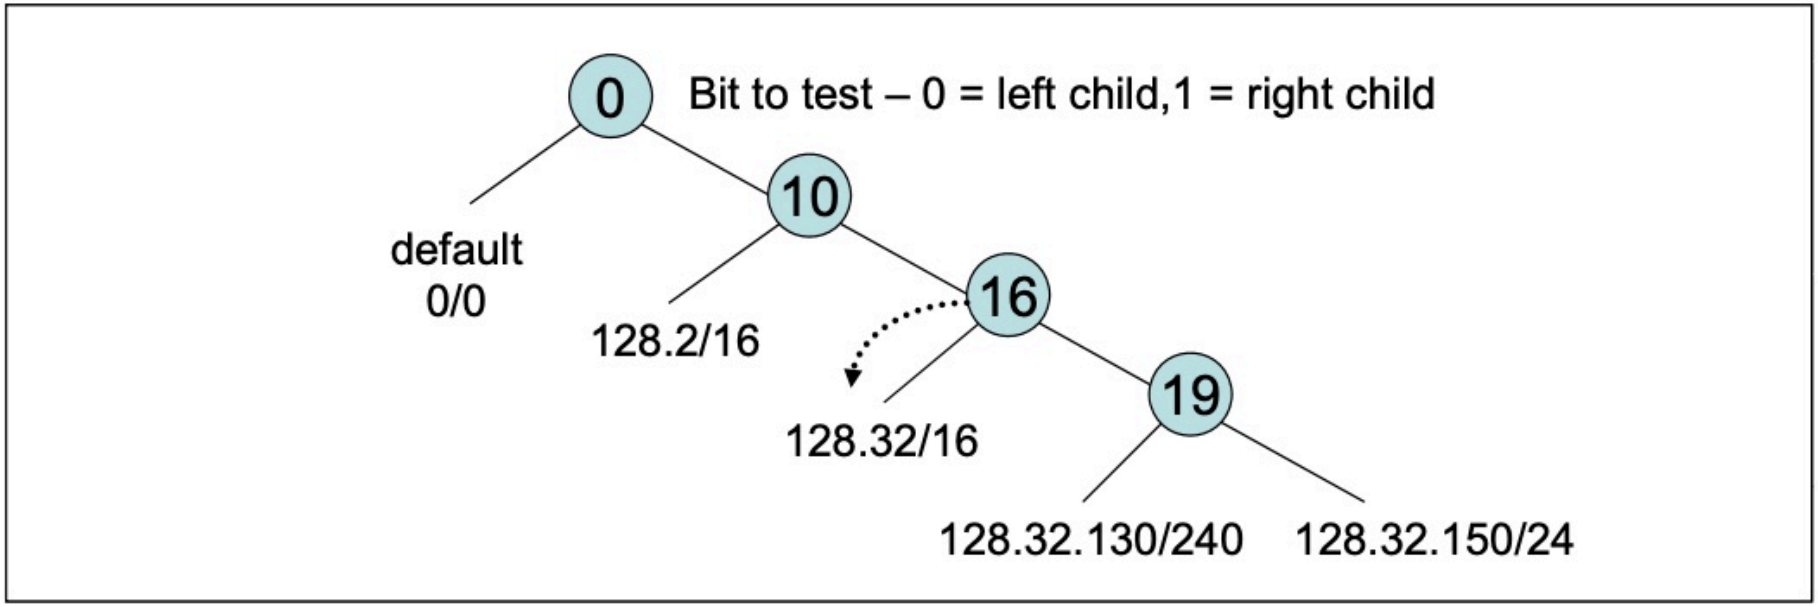

Address Lookup: Direct Trie

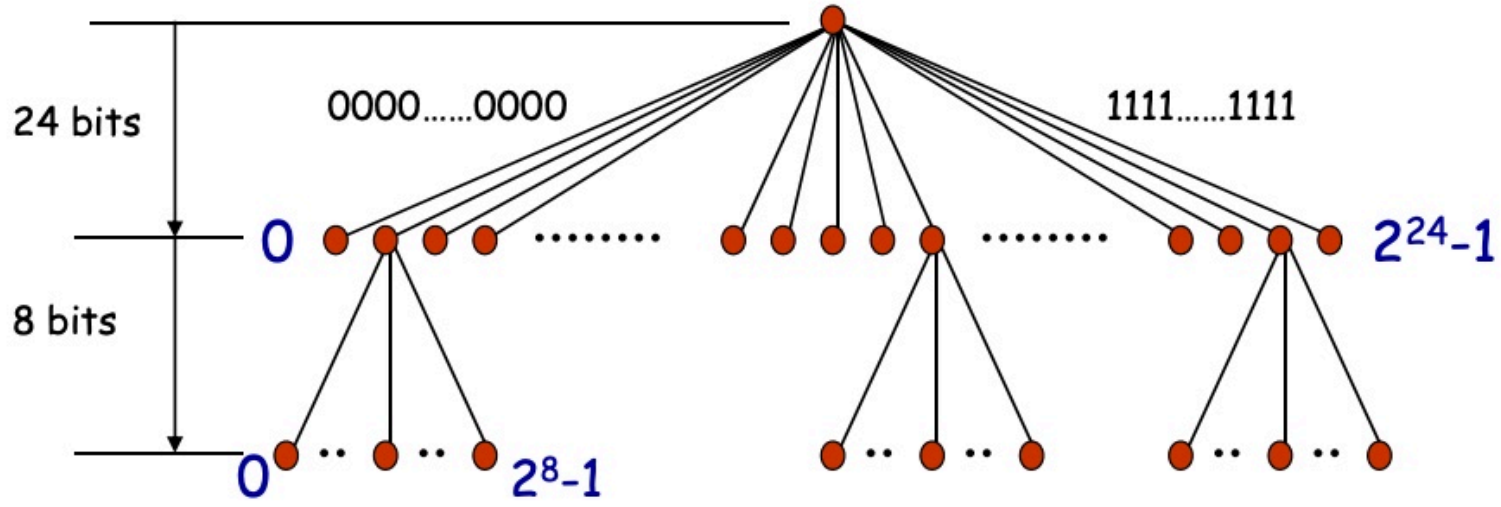

- When pipelined, one lookup per memory access
- **Inefficient use of memory**

Hardware for LPM lookup

- Content-Address Memory (CAM)
  - Input: tag (address)
  - Output: value (port
  - Exact match, but O(1) in hardware
- Ternary CAM
  - Can have wildcards: 0, 1, *
  - “value” memory cell and “mask” (care / don’t care) cell
- LPM via TCAM
  - In parallel, search all prefixes for all matches
  - Then choose longest match
    - Trick: choose first match, but already sorted by prefix length

Example: LPM with a TCAM

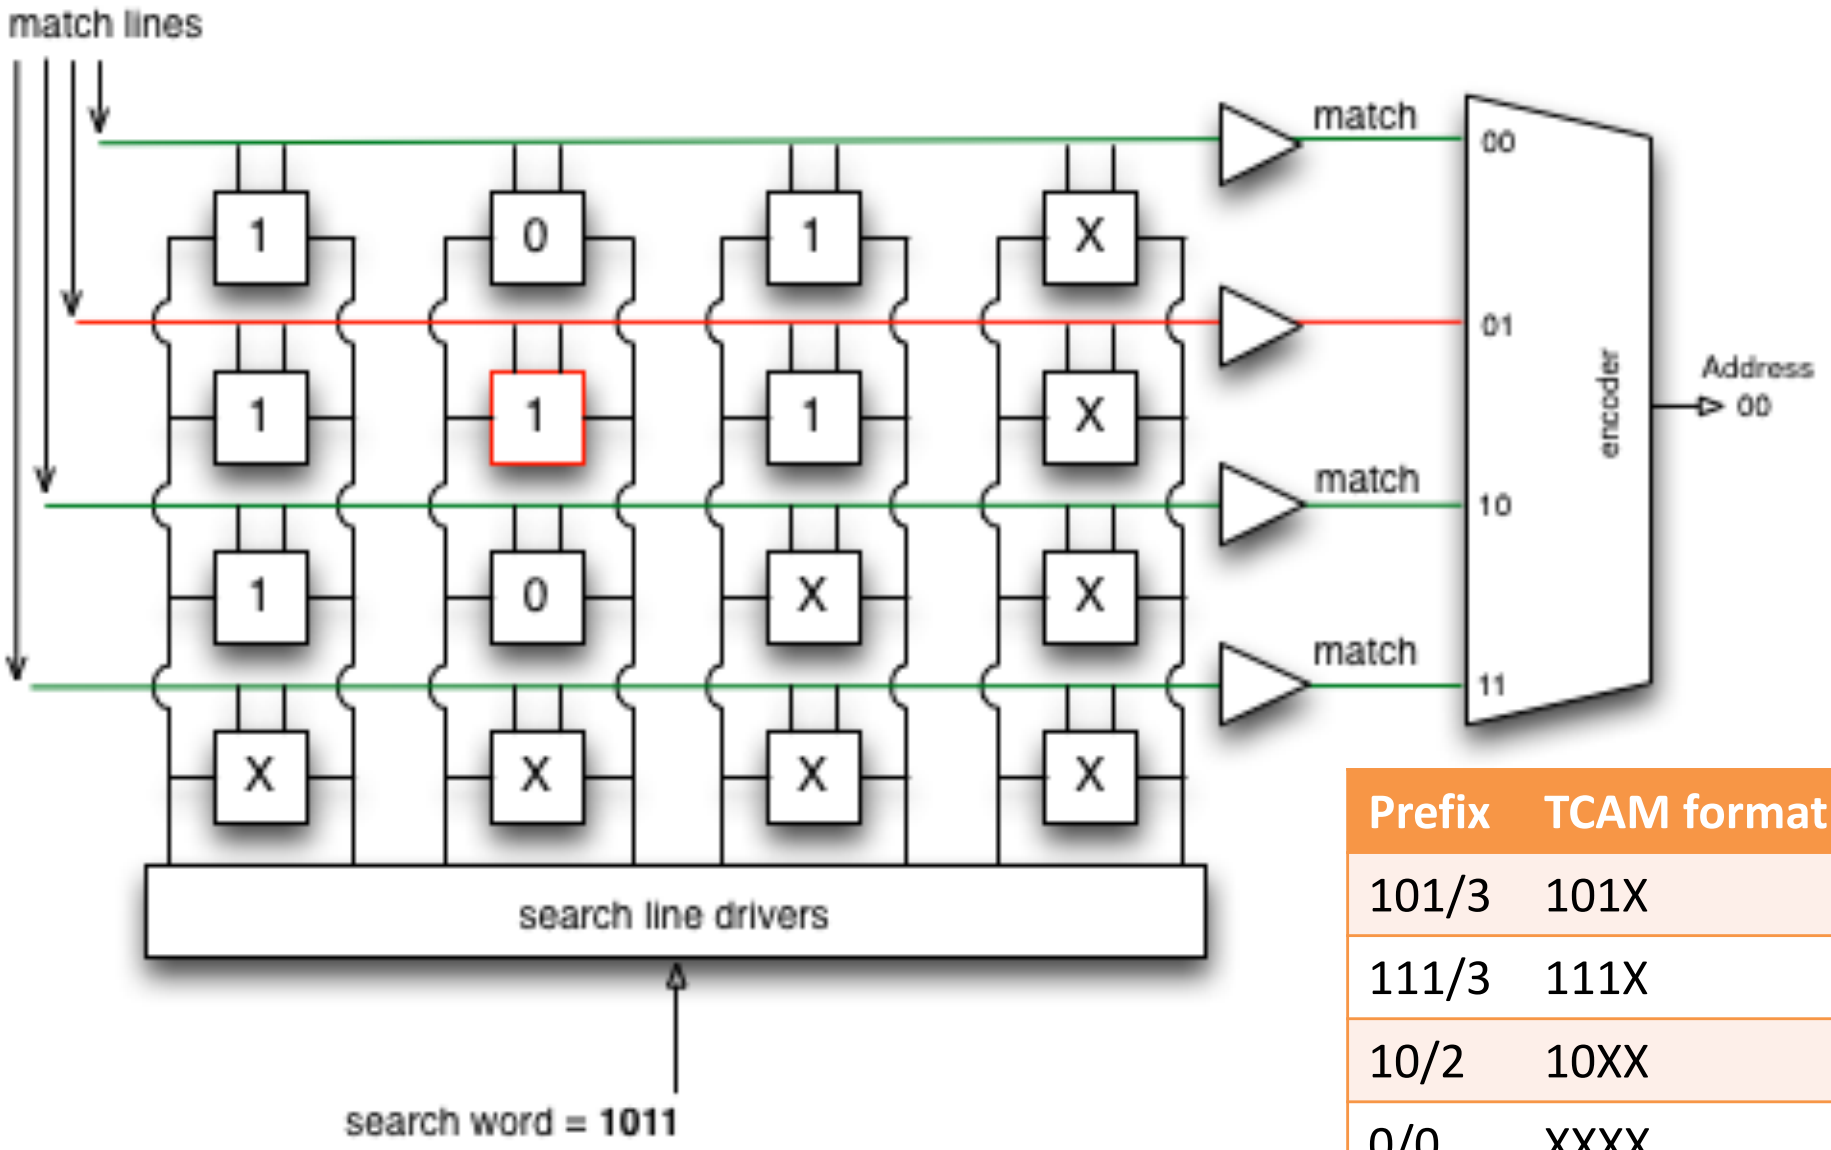

Prefix	TCAM format
101/3	101X
111/3	111X
10/2	10XX
0/0	XXXX

[thenetworksherpa.com/tcam-in-the-forwarding-engine]

Decision: Forwarding table per line card

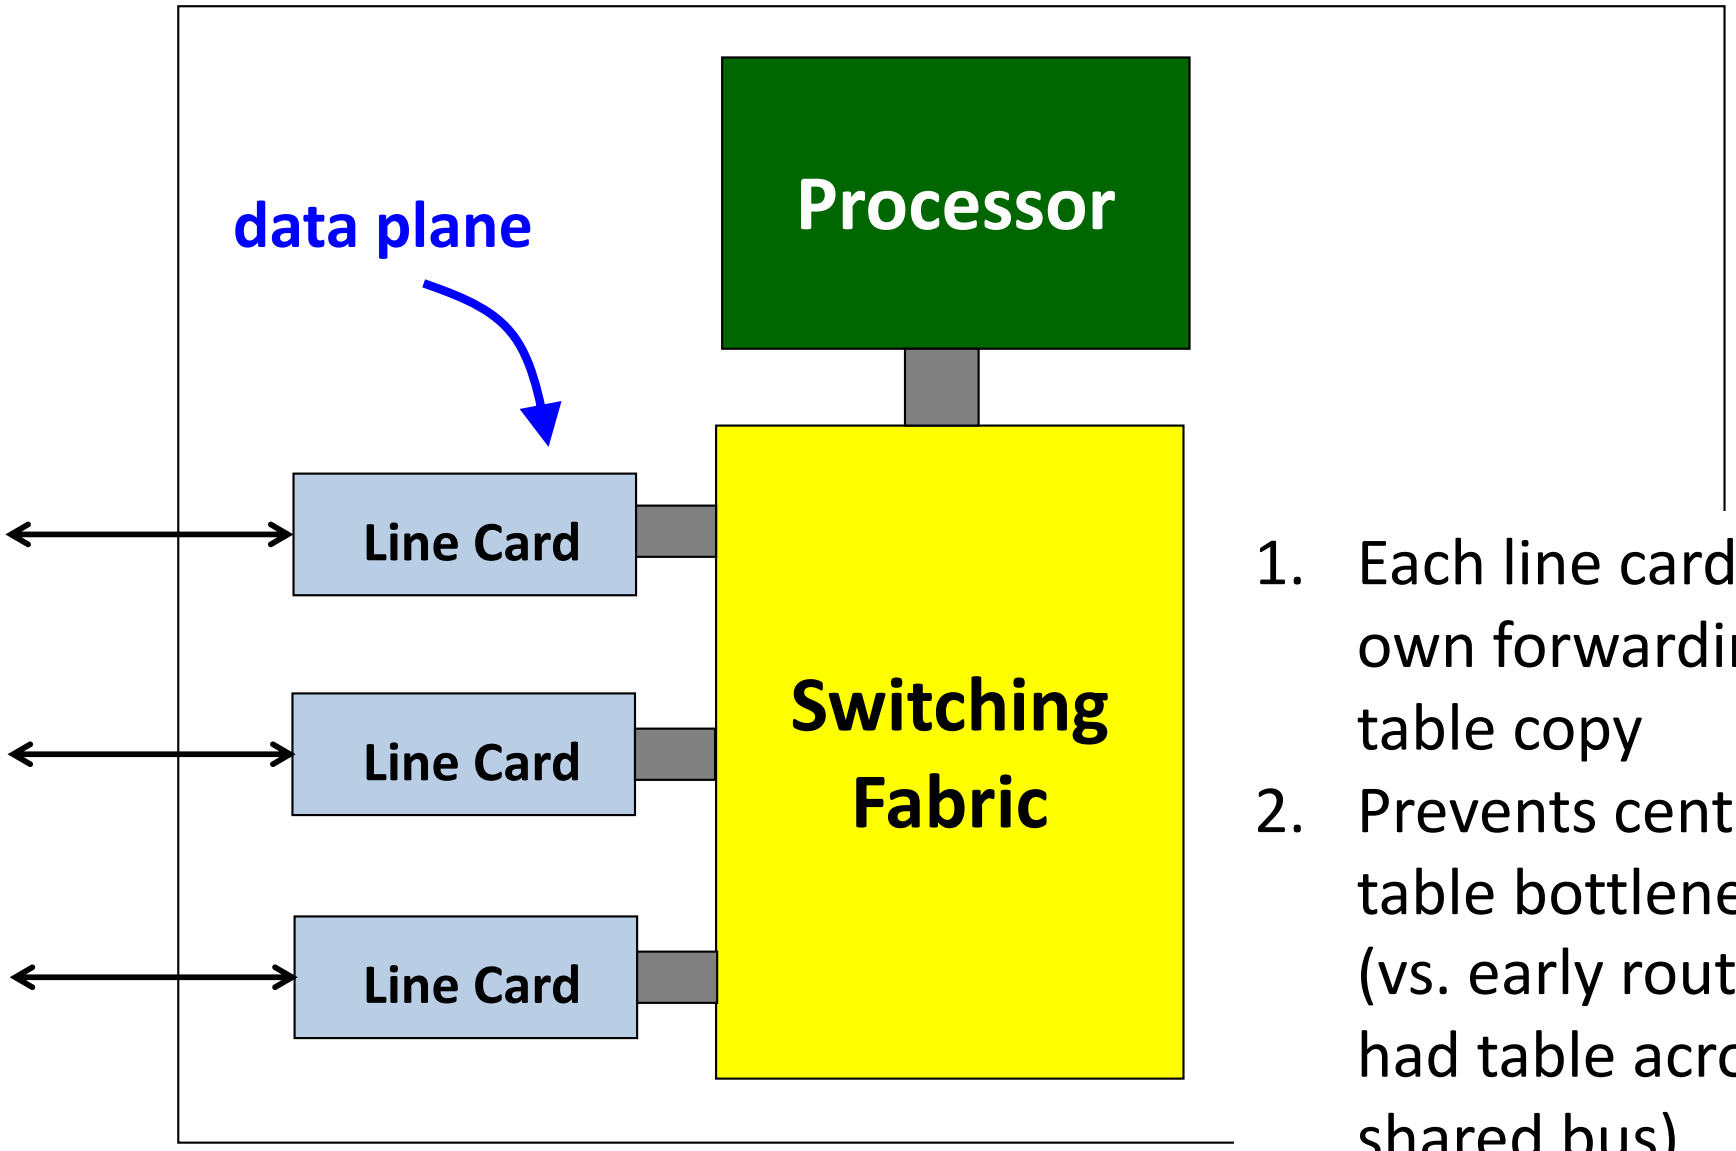

1. Each line card has own forwarding table copy
2. Prevents central table bottleneck (vs. early routers had table across shared bus)

Decision: Crossbar switch

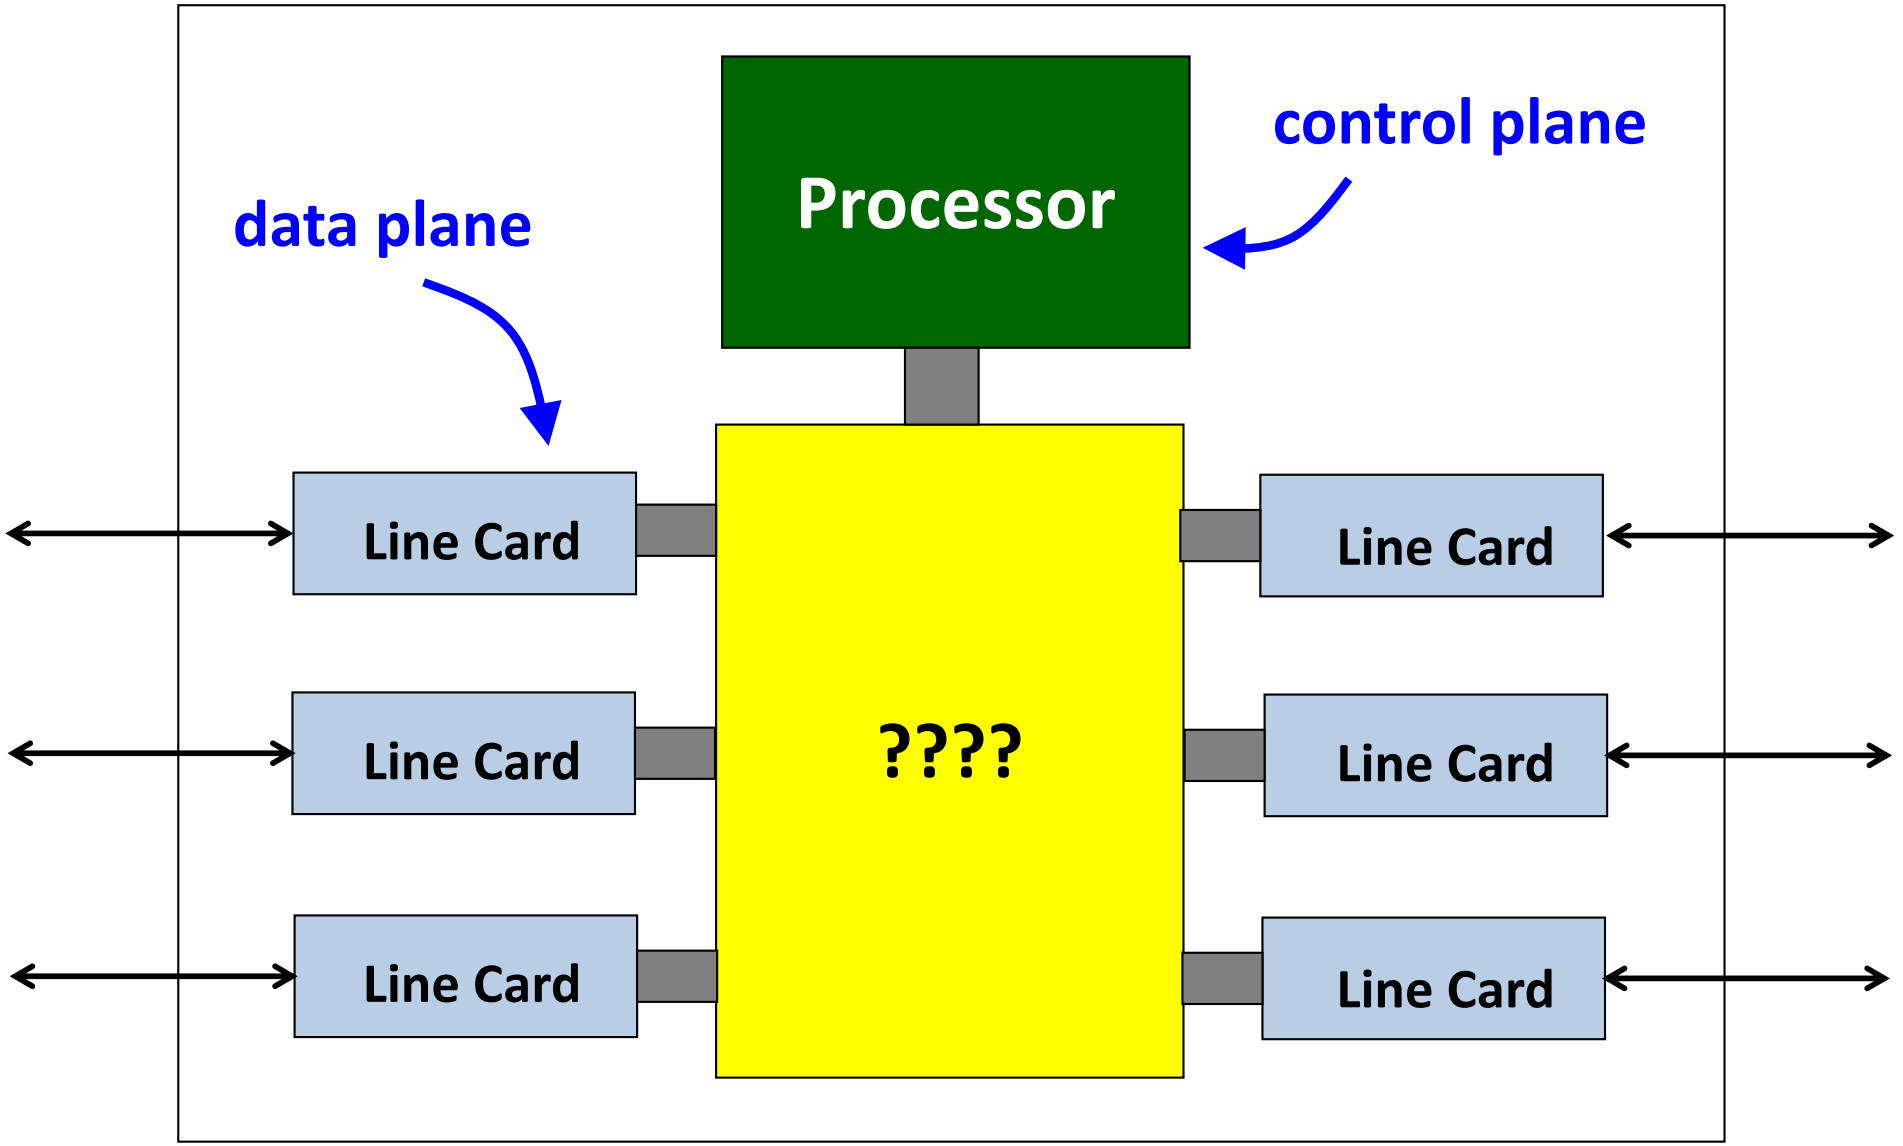

Decision: Crossbar switch

- Shared bus
  - Only one input can speak to one output at a time
- Crossbar switch / switched backplane
  - Input / output pairs that don't compete can send in same timeslot

Crossbar switching

- Every input port has connection to every output port
- In each timeslot, each input connected to zero or more outputs

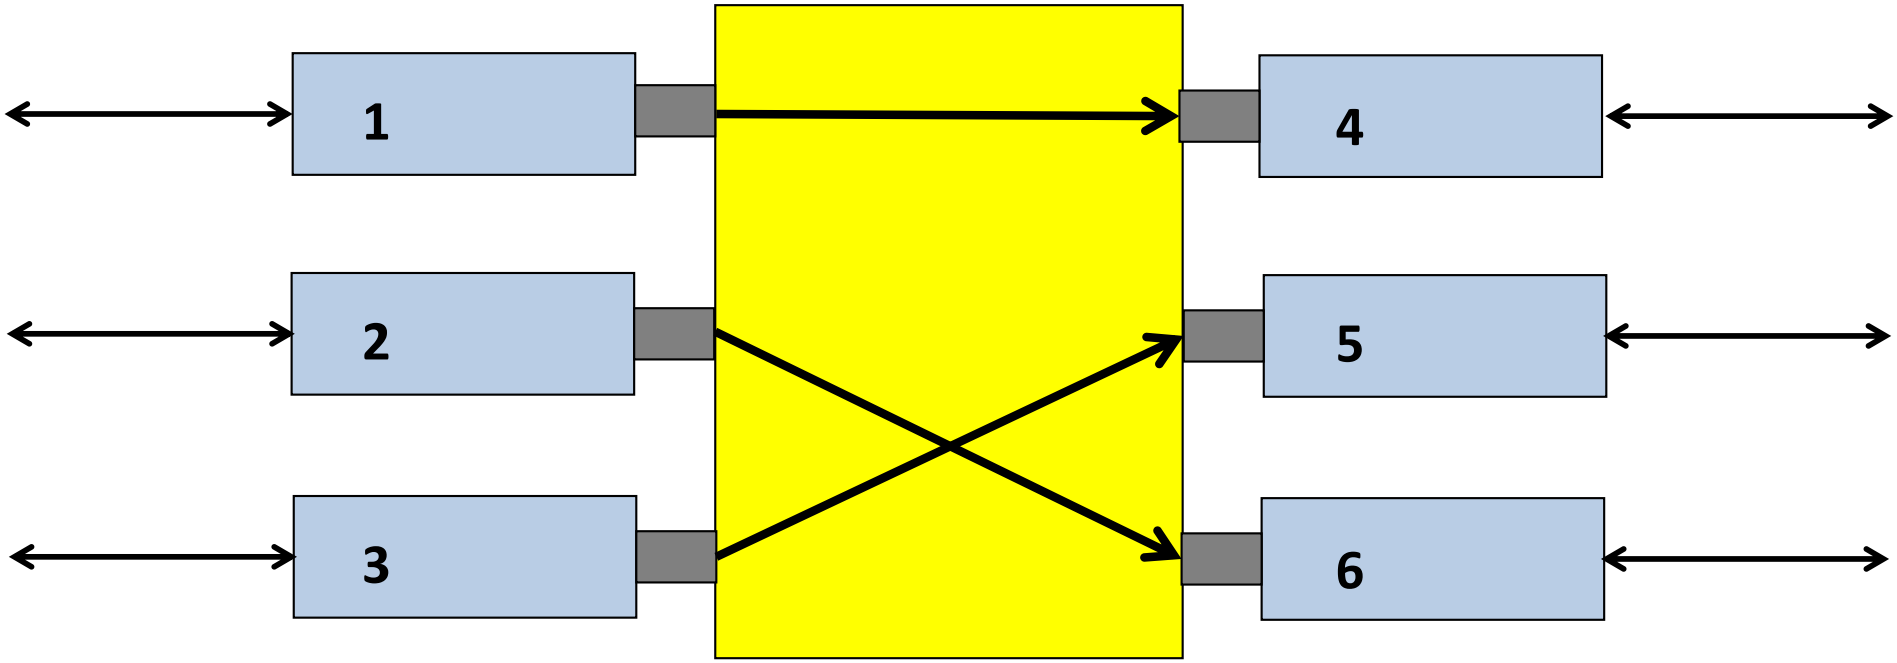

Crossbar switching

- Every input port has connection to every output port
- In each timeslot, each input connected to zero or more outputs

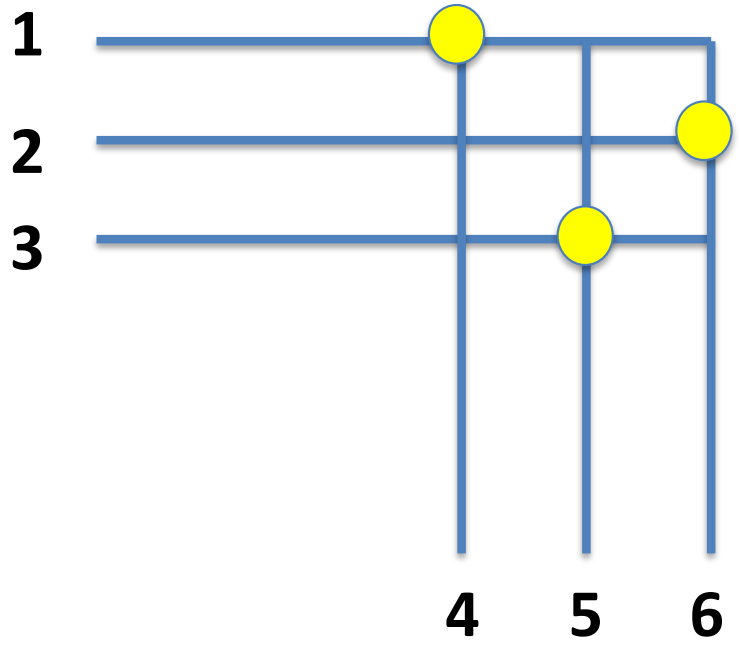

- Good parallelism
- Needs scheduling

Everything gets complicated…

- **Problem: Head-of-line blocking**
  - The packet in front of queue blocks packets behind it from being processed
  - Say first packet at input 1 wants to go to output 5; second packet at 1 that wants 4 is still blocked

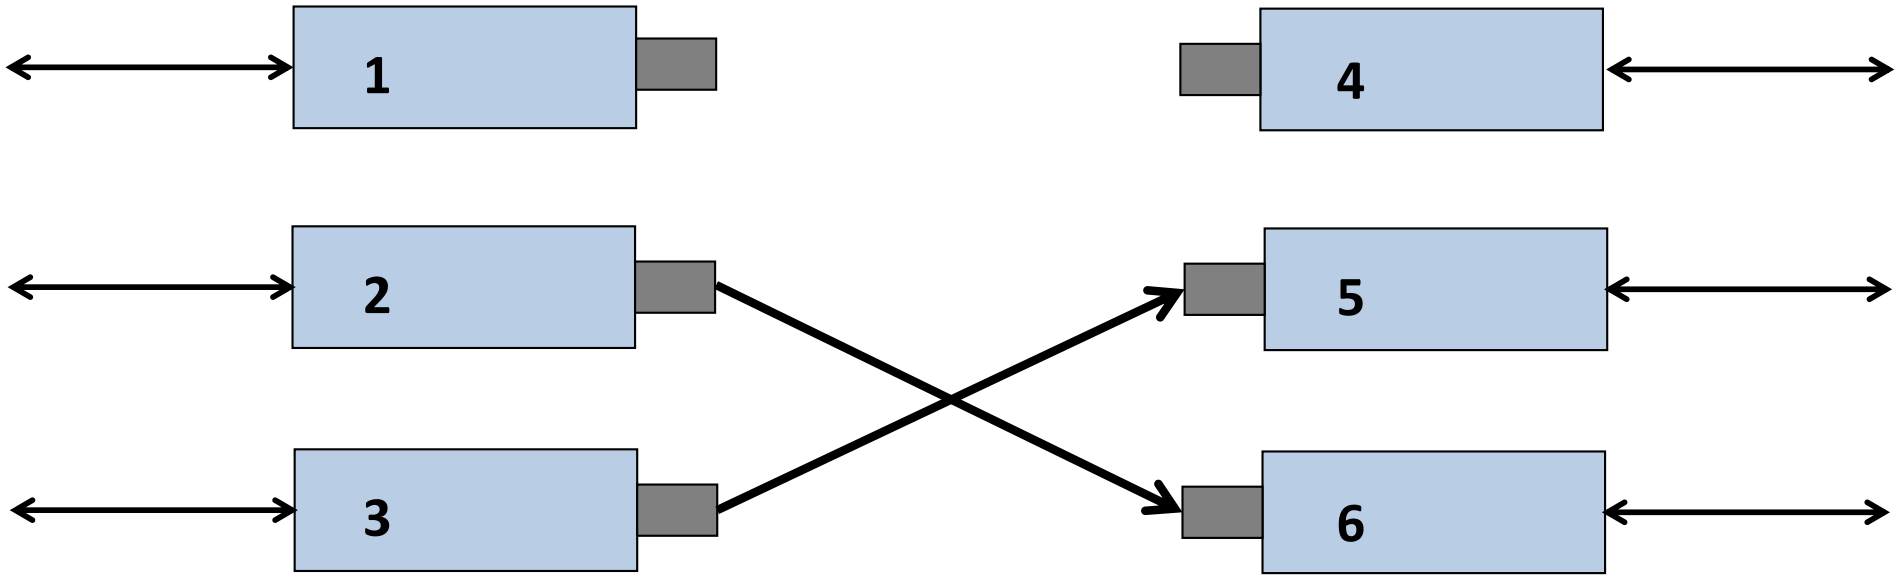

Everything gets complicated…

- **Solution: *Virtual output queues***
  - One queue at input, **per output port** (for all inputs)
  - So **avoids head-of-line blocking** during crossbar scheduling

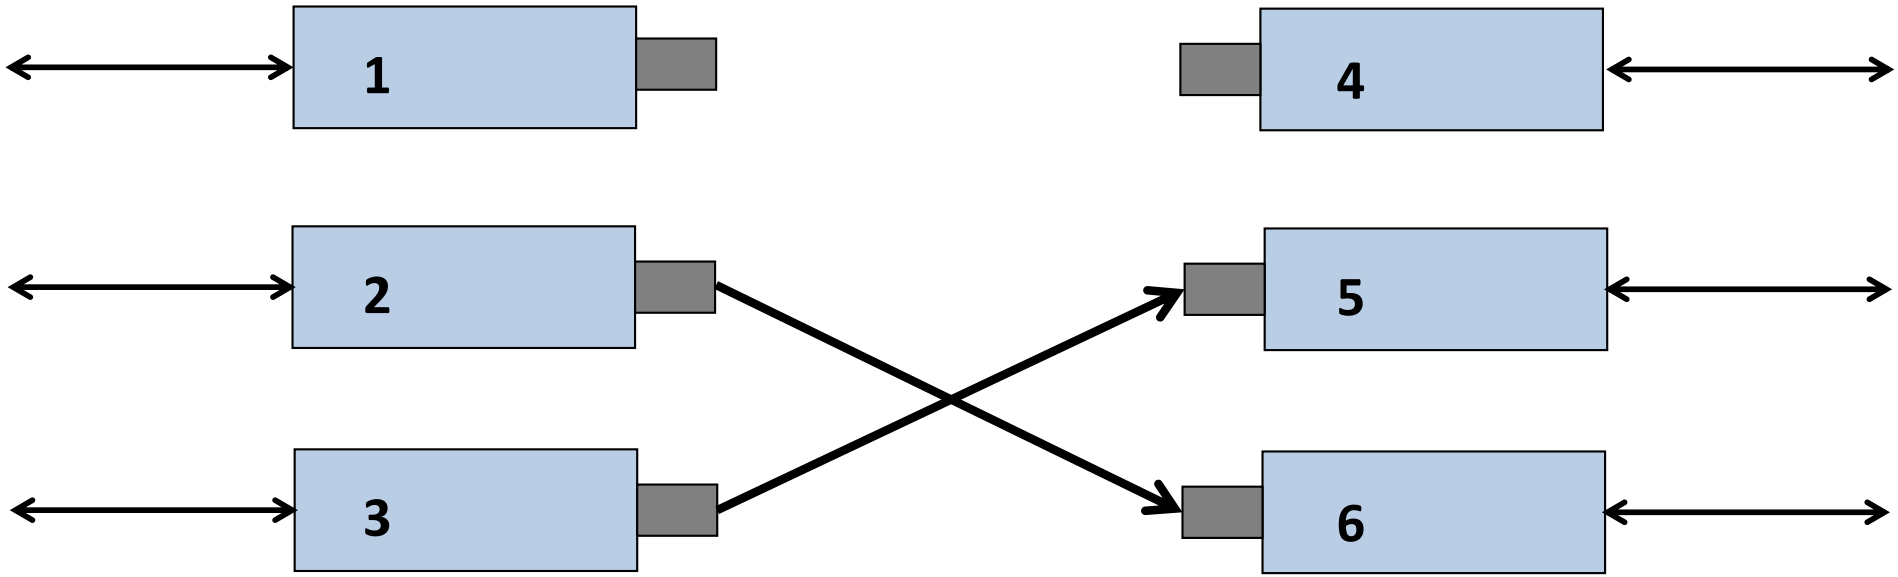

Cisco 8000 Series Routers

- Up to 648 400 GbE
- 260 Tbps backplane

Collision and broadcast domains

- Hub

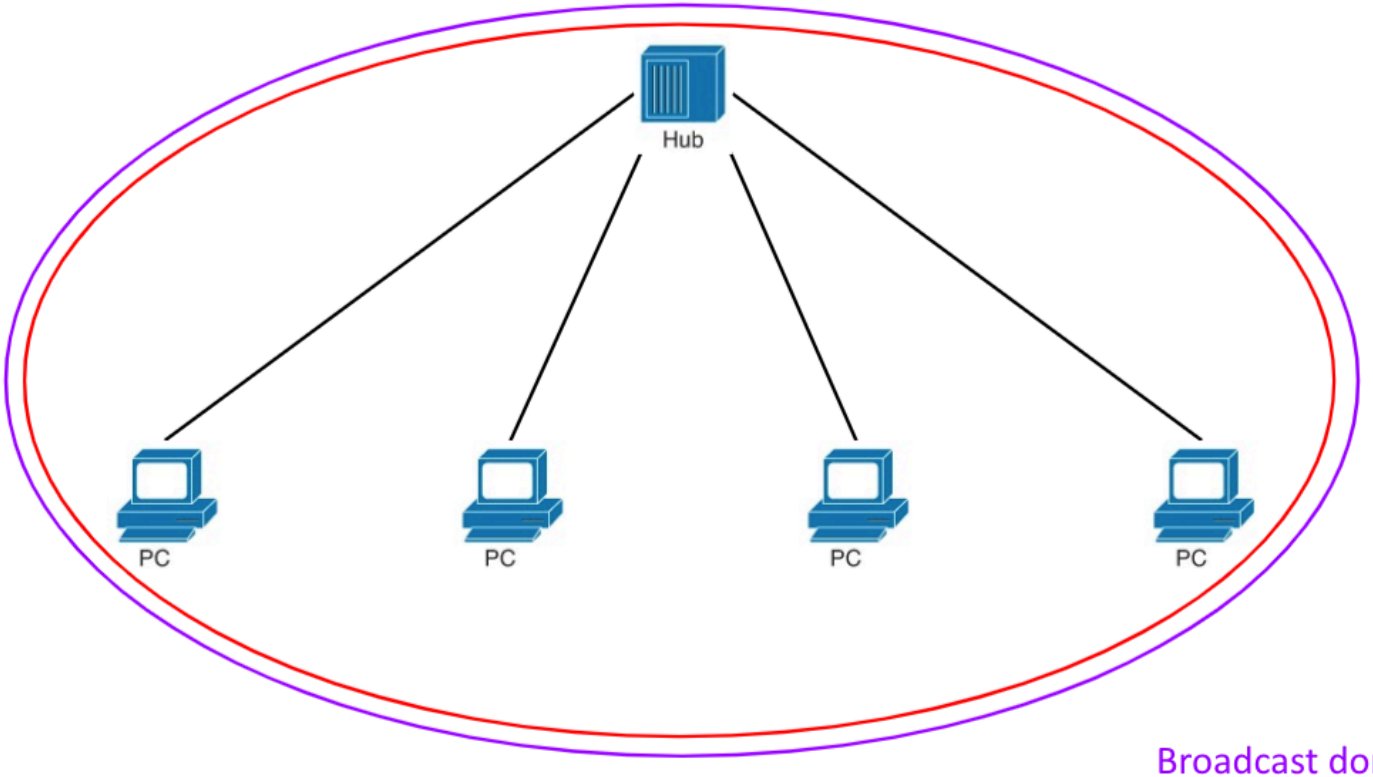

Collision and broadcast domains

- Switch

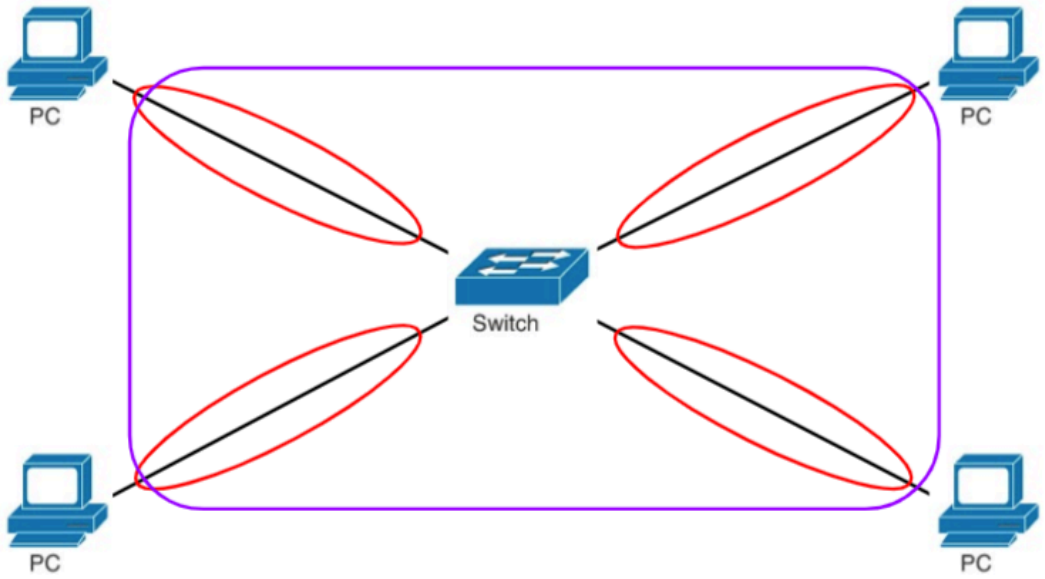

Broadcast domain

Collision domain

Collision and broadcast domains

- Switch and router

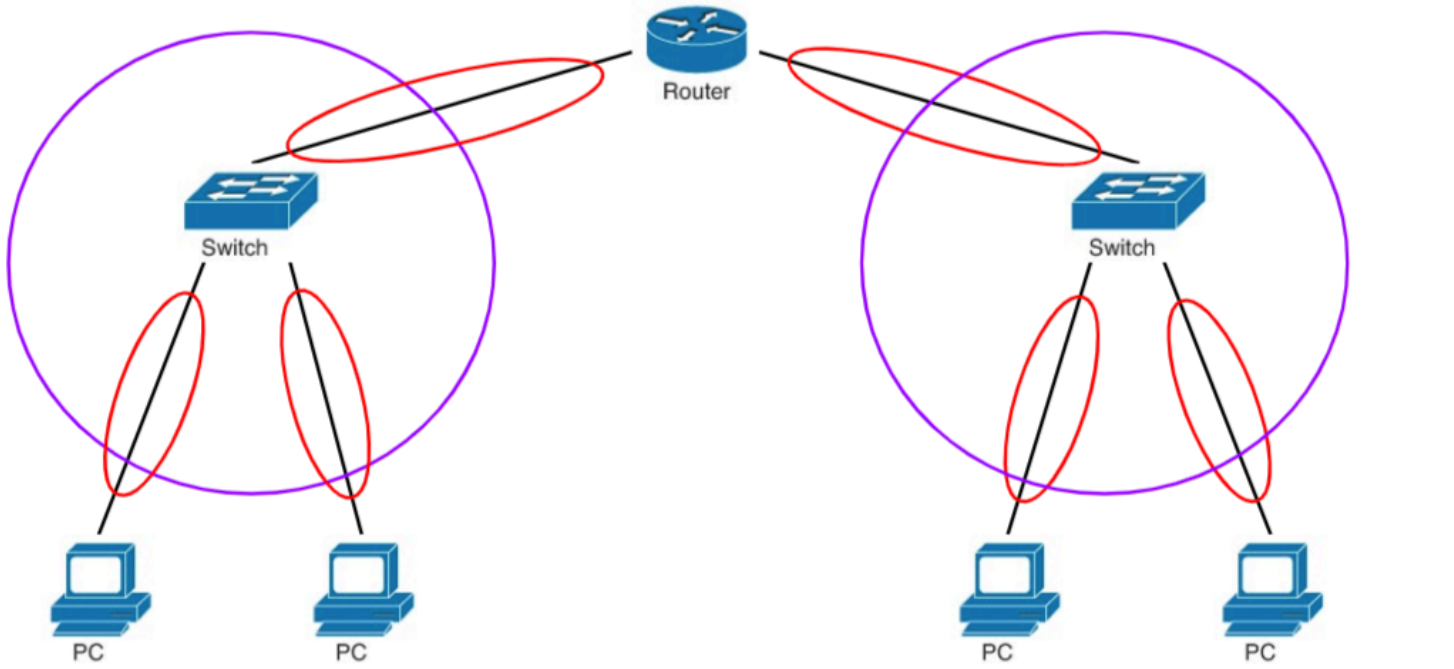

Collision and broadcast domains

- Hub, switch, and router

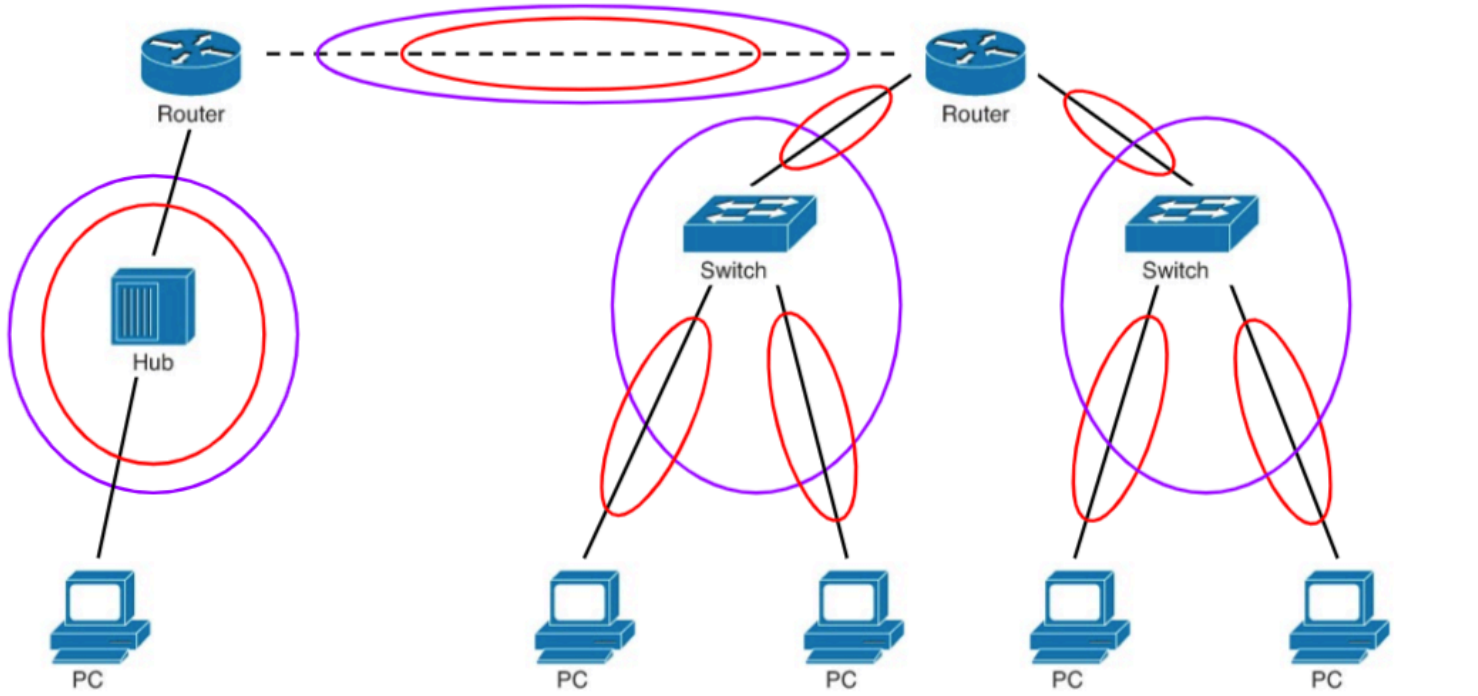

Conclusions

- Physical devices sharing L2 & L3 networks have many common features
  - Forward table lookups
  - Queueing and backplane switching
  - Fast vs. slow paths
    - Switches and routers separate routing decisions (control plane) from forwarding actions (data plane)
- High speed necessitates innovation
  - Specialized hardware
  - Software algorithms